



Data Fabric Solutions

Data Fabric Solution for Cloud Backup Workflow Guide

Using ONTAP[®], SnapCenter[®], and AltaVault[®]

January 2017 | 215-11811_A0
doccomments@netapp.com

Contents

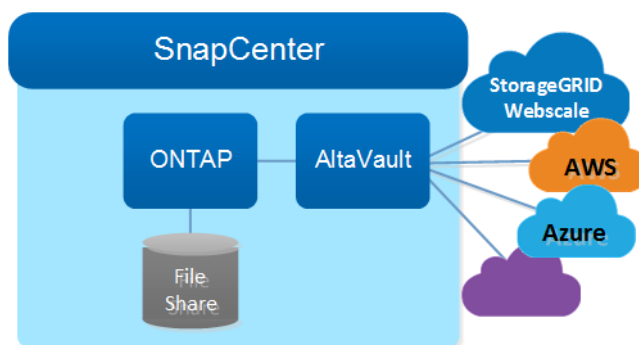
Data Fabric Solution for Cloud Backup overview	5
Data Fabric Solution for Cloud Backup components	6
Data Fabric Solution for Cloud Backup features	6
What you can do with Data Fabric Solution for Cloud Backup	8
Preparing for the Data Fabric Solution for Cloud Backup deployment	9
ONTAP requirements	9
SnapCenter host and file catalog requirements	10
AltaVault requirements	11
Supported browsers	12
Licensing requirements	12
Connection and port requirements	15
Credential requirements	15
Installation and configuration workflow with SnapCenter data protection management	16
Installation and configuration workflow without SnapCenter data protection management	17
Installing the Data Fabric Solution for Cloud Backup	18
Upgrading ONTAP	18
Establishing cluster peering for Data Fabric Solution for Cloud Backup	19
Configuring intercluster LIFs to use dedicated intercluster ports	19
Adding an ONTAP license	22
Installing SnapCenter for Data Fabric Solution for Cloud Backup	23
Ensuring the Linux server is accessible for NAS file services file catalog ...	23
Configuring sudo privileges for non-root user	23
Installing the SnapCenter Server for Data Fabric Solution for Cloud Backup	24
Backup	24
Installing AltaVault	26
Configuring AltaVault for Data Fabric Solution for Cloud Backup	27
Logging in to AltaVault	27
Enabling the SnapMirror service and whitelist IP connections using the AltaVault GUI	28
Enabling SnapMirror long-term retention mode in AltaVault	29
Enabling SnapCenter access to AltaVault	30
Adding a SnapCenter role-based user account to enable SnapCenter storage connections to an AltaVault appliance	30
Starting the AltaVault Storage Optimization Service	31
Configuring SnapMirror and enabling SnapCenter access using the AltaVault CLI	31
Configuring SnapCenter for Data Fabric Solution for Cloud Backup	34

Logging in to SnapCenter	34
Configuring the Infrastructure Admin role for AltaVault access	35
Assigning users access to assets	36
Setting up storage system connections	37
Setting up your Run As credentials	39
Configuring CIFS servers for Data Fabric Solution for Cloud Backup	41
Protecting the SnapCenter file catalog and the repository	41
Preparing the PowerShell environment	42
Protecting the file catalog and SnapCenter repository with PowerShell cmdlets	43
Backing up NAS file services to the cloud and restoring them with SnapCenter management	45
Using policies to govern backing up NAS file services	45
Backing up NAS file services to the cloud	47
Backing up NAS file services to the cloud using SnapCenter and AltaVault	48
Monitoring data protection progress in the SnapCenter Activity pane	49
Monitoring backup operations from the Jobs page	50
Verifying backup to the cloud completion on AltaVault	50
Restoring NAS file data from the cloud	53
Searching in the file catalog for a specific file to restore using keywords and filters	54
Restoring NAS file data from the cloud using SnapCenter and AltaVault	56
Monitoring restore operations from the Jobs page	57
Managing SnapMirror shares using the AltaVault CLI	57
Backing up NAS file services to the cloud and restoring them using ONTAP commands	59
Backing up NAS file services to the cloud using ONTAP commands	59
Performing a baseline backup from ONTAP to the cloud using the ONTAP CLI	59
Performing an incremental backup from ONTAP to the cloud using the ONTAP CLI	62
Monitoring backup progress using the ONTAP CLI	63
Restoring NAS file services from the cloud using ONTAP commands	64
Restoring a full backup from the cloud using the ONTAP CLI	64
Restoring a single file from the cloud to ONTAP using the ONTAP CLI	67
Viewing data protection reports for NAS file services	69
Where to go next	70
Copyright information	71
Trademark information	72
How to send comments about documentation and receive update notifications	73
Index	74

Data Fabric Solution for Cloud Backup overview

With the Data Fabric Solution for Cloud Backup, storage administrators and IT generalists can perform Snapshot copy-based data protection operations for Windows and UNIX file shares to and from public or private cloud object stores. Delivering high service levels with fast Snapshot copies and efficient array-based replication, the solution includes NetApp ONTAP, SnapCenter data management software, and AltaVault cloud-integrated storage software.

ONTAP provides storage management and initiation and management of protection operations using a command-line interface (CLI). SnapCenter provides support for protection operations initiation and management. AltaVault provides storage-optimized space in cloud storage providers for Snapshot copies of file share data from ONTAP.

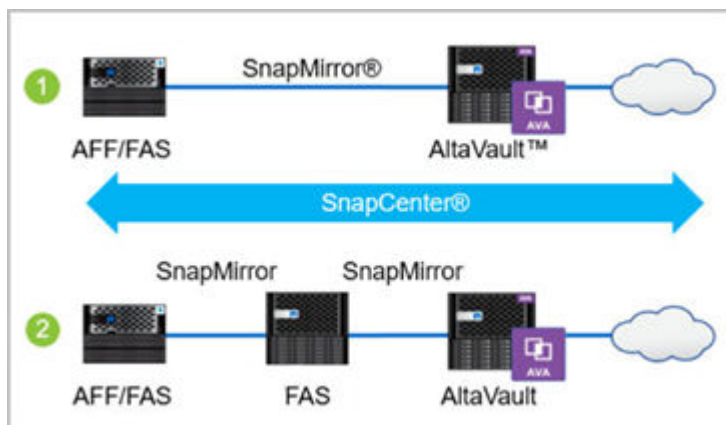


The Data Fabric Solution for Cloud Backup using SnapMirror technology enables administrators to back up NFS and SMB NAS file services (NFS v4, SMB3) and home directories to the cloud using an AltaVault cloud provider connection.

The ONTAP CLI can be used for additional operations not supported by the SnapCenter graphical interface. For example, using the ONTAP CLI, you can also back up ONTAP LUNs.

Flexible configurations are available for All Flash FAS (AFF) or FAS in these topologies:

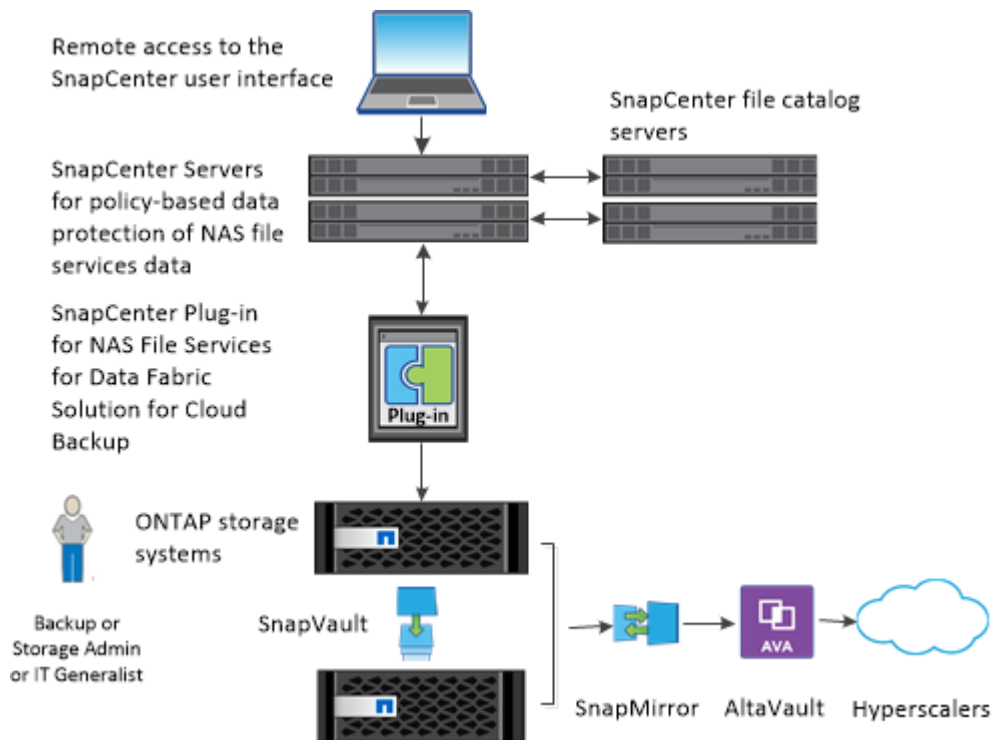
- Primary storage (AFF or FAS) to AltaVault
- Primary storage (AFF or FAS) to secondary storage (FAS) to AltaVault



Data Fabric Solution for Cloud Backup components

Data Fabric Solution for Cloud Backup involves multiple products that together provide data protection of NAS file services to public or private cloud providers.

- ONTAP 9.1 with either a SnapMirror or SnapVault license
- SnapCenter 2.0 and the SnapCenter Plug-in for NAS File Services
- AltaVault 4.3 (physical or virtual)



ONTAP provides the foundation for data protection operations of NAS file services. Using ONTAP CLI, you can initiate, manage, and monitor data protection operations.

SnapCenter provides support for the management of policies, schedules, and backup and restore operations for NAS file services.

AltaVault provides optimized storage for backup and restore operations, backed by both public and private cloud providers, including NetApp StorageGRID Webscale, Amazon Web Services, Microsoft Azure, Google Cloud Storage, and IBM SoftLayer. When an administrator creates a SnapMirror relationship in ONTAP with AltaVault, a share is automatically created in AltaVault. When you use SnapCenter to initiate backups to AltaVault, the relationship is automatically created.

Each share is associated with one ONTAP FlexVol volume. AltaVault supports up to 500 SnapMirror shares.

Data Fabric Solution for Cloud Backup features

Each component in the Data Fabric Solution for Cloud Backup contributes to the completion of cloud-integrated data protection services for NAS file services.

The Data Fabric Solution for Cloud Backup includes the following key features using ONTAP, SnapCenter, and AltaVault in combination:

- Ability to transport NAS file services data on primary or secondary storage systems to private or public cloud service providers
- Data protection of ONTAP FlexVol volumes to the cloud via AltaVault
- Choice of All Flash FAS (AFF) or FAS hybrid storage for primary storage and FAS on optional secondary storage backup
- Use of Snapshot copy technology and SnapMirror data protocol, which enables fast and efficient data transfer between ONTAP and AltaVault without having to stream data through a backup server and then to the AltaVault appliance
- Cost-effective tape replacement solution

ONTAP includes these key features:

- Data management across flash and disk
 - Data protection of file workloads for this solution
 - Ability to leverage the architecture you want: NetApp engineered systems, software-defined storage, and the cloud
 - Deployment of NAS workloads for this solution on a unified storage architecture
 - Storage footprint and cost reduction with inline data reduction
- When you use the Data Fabric Solution for Cloud Backup, data reduction occurs in AltaVault.

Note: This solution does not support ONTAP Select or ONTAP Cloud.

SnapCenter enables self-service data management with role-based access control. With support for NAS file services, SnapCenter provides these additional features:

- Single management interface to manage NAS file services along with other applications and databases in the enterprise
- Access governed by role-based access control
- Preconfigured backup policies that simplify management for NAS file services
- End-to-end protection automation with policy-based management and secondary and tertiary relationship configuration
- Scalable, highly available file catalog for quick location and recovery of single files
- Search or browse option to identify a file for restore quickly
- Visual representation of backup copies that lets you quickly identify backup to restore

AltaVault include these key features:

- Ability to transport data to both public and private cloud providers
- Note:** Data Fabric Solution for Cloud Backup does not currently support AWS Glacier.
- Accelerated data protection with the deduplication and compression of data volumes at ingest
Data transmitted from ONTAP to AltaVault is not deduplicated or compressed. All data stored on AltaVault is deduplicated and compressed upon ingest and replicated over the Internet to cloud storage for efficient WAN transmission.
 - Data encryption applied at ingest and securely transmitted to the cloud storage provider, reducing security and compliance risks locally and in the cloud
 - Flexible deployment and scale with physical or virtual environments

What you can do with Data Fabric Solution for Cloud Backup

Using ONTAP, SnapMirror technology, SnapCenter, and AltaVault, administrators can protect their NAS file services on FlexVol volumes in public or private clouds.

With the Data Fabric Solution for Cloud Backup and ONTAP, you can accomplish the following:

- Perform a single baseline Snapshot copy-based backup and then block-level incremental backups forever from ONTAP to AltaVault and then to public or private cloud providers.
- Perform a single file restore of data from AltaVault to ONTAP using SnapCenter. With the ONTAP CLI, you can perform a full volume restore or a single file restore.
The ONTAP CLI can be used for additional operations not supported by the SnapCenter graphical interface. For example, using the ONTAP CLI, you can also back up ONTAP LUNs.

Using SnapCenter, you can accomplish the following:

- Back up NAS file services (on FlexVol volumes) from flash to disk to AltaVault and to the cloud.
- Apply predefined or custom policies for your backup and retention requirements.
- Restore individual files within file shares to the original location on the same volume.
Note: You might want to rename the current file if it exists, so it is not overridden.
- Search for data to restore in one of two ways. Either use a powerful keyword search in the SnapCenter file catalog and apply filters to find specific files quickly among millions of files or, from a selected backup, find and restore a file
- Manage backup schedules and policies.

Using AltaVault, you can complete the cloud backup protection by performing these tasks:

- Perform efficient and secure backup to most clouds, public or private.
- Reduce I/O traffic to the cloud, saving ingest costs and in-cloud storage capacity with inline deduplication and compression.
- Accomplish long-term retention in the cloud.
- Connect multiple FAS systems to a single AltaVault. Using ONTAP CLI, you can manage this fan-in configuration relationship.

Preparing for the Data Fabric Solution for Cloud Backup deployment

Before you install and set up Data Fabric Solution for Cloud Backup, you must prepare your environment and understand the information that you need for installation and setup.

ONTAP requirements

Before you begin the Data Fabric Solution for Cloud Backup deployment, you should be familiar with the ONTAP requirements.

Item	Minimum
ONTAP version	ONTAP 9.1 or greater
Models	No restrictions on models. For storage system requirements for ONTAP 9.1, see the Hardware Universe. NetApp Hardware Universe
Storage type	Primary. Any system that runs ONTAP 9.1 or greater. FAS and All Flash FAS (AFF) only.
Storage amount	No minimum requirements.
Intercluster LIFs	Required.
Cluster peering and SVM peering	Cluster peering requirements differ based on the type of backup: <ul style="list-style-type: none"> For backups from ONTAP to AltaVault, cluster peering is not required. For backups from ONTAP to ONTAP to AltaVault, cluster peering and storage virtual machine (SVM) peering are required. The process for configuring cluster and SVM peering differs depending on whether peering has already been configured: <ul style="list-style-type: none"> If cluster peers are not already configured, you must configure both cluster and SVM peering manually. If cluster peers are already configured, you do not need to configure SVM peering; it is done automatically.

Primary FAS or All Flash FAS (AFF) systems directly connected to AltaVault require ONTAP 9.1 or greater. If you want to recover using AltaVault directly to the primary FAS or AFF, the primary system requires ONTAP 9.1 or greater.

In a cascade configuration where SnapMirror relationships exist such as from ONTAP to ONTAP, the secondary system requires ONTAP 9.1 or greater, but the primary system could use ONTAP 8.3.1 or greater. The Data Fabric Solution for Cloud Backup supports cascade configurations as well where AltaVault becomes the endpoint. The restore path in this configuration requires two steps: restoring from AltaVault to secondary FAS and then from secondary FAS to primary FAS or AFF.

For the latest information, see the Interoperability Matrix.

Related information

[NetApp Interoperability Matrix Tool](#)

SnapCenter host and file catalog requirements

Before you begin the Data Fabric Solution for Cloud Backup deployment, you should be familiar with SnapCenter host and file catalog requirements.

SnapCenter Server requirements

The SnapCenter Server should meet the following minimum requirements. The SnapCenter Server should have a dedicated resource; it should not be shared.

All supported Windows operating systems are 64-bit only.

The Data Fabric Solution for Cloud Backup requires MySQL Server on new SnapCenter installations only. Upgrades from existing SnapCenter installations are not supported.

Item	Minimum
Version	2.0
CPU	4 CPUs
Memory	16 GB
Operating System	<p>The following are supported:</p> <ul style="list-style-type: none"> Windows Server 2016 Windows Server 2012 DCE, R2, Se, R2 SE <p>Microsoft has identified an HTTP security vulnerability in their Windows operating systems that requires the installation of KB3042553-x64.</p> <p>https://support.microsoft.com/en-us/kb/3042553</p>
SnapCenter storage amount	40 GB for both the software and the repository
SnapCenter metadata repository	<p>MySQL Server 5.7 or greater</p> <p>The installation installs the SnapCenter repository on drive C of the host. After installation, you might want to move the repository using the <code>protect-SMRepository</code> PowerShell cmdlet.</p>
Third-party libraries required on the SnapCenter Server	<ul style="list-style-type: none"> Java 1.8.0.71 or later (Oracle or OpenJDK), 64-bit version. Must be installed before installing SnapCenter. File archiver utility p7zip-9.20.1-8.1.1.x86_64.rpm or exe. Included in installation.

For the latest information, see the Interoperability Matrix.

SnapCenter file catalog requirements

One or more servers or VMs should be dedicated to the file catalog.

We recommend two separate nodes for the SnapCenter file catalog, with each node having enough disk space for indexing, depending on how much data will be indexed. We also recommend a replica for each partition shard for file catalog high availability and disaster recovery.

Item	Description
CPU	8 cores
Memory	24 GB RAM
Operating system	Red Hat Enterprise Linux (RHEL) 7.0, 7.1, and 7.2
Dedicated or shared server	Dedicated VM
Supported Hypervisors	VMware, Hyper-v
Java requirements	JRE 1.8
Linux host permissions	Root access is not required. You must have enabled the password-based SSH connections for the root or non-root user. You must configure the sudo privileges for the non-root user.

We recommend the following file catalog storage depending on the number of files to be indexed:

Total number of files	Size	Disk capacity per node	CPU per node	Physical RAM per node	Shards	Replica per shard	Total number of machines
100 million	50 GB	250 GB	8 cores	12 GB	1	1	2
200 million	100 GB	250 GB	8 cores	24 GB	1	1	2
400 million	250 GB	250 GB	8 cores	48 GB	1	1	2
1 billion	500 GB	500 GB	8 cores	120 GB	1	1	2

A LUN must be provisioned for storing the file catalog data. During installation, you can specify the path to the preconfigured, provisioned LUN.

AltaVault requirements

Before you begin the Data Fabric Solution for Cloud Backup deployment, you should be familiar with the AltaVault requirements.

Item	Minimum
Models	AltaVault physical models: AVA-400, AVA-800 AltaVault virtual models: AVA-v2, AVA-v8, AVA-v16, AVA-v32 AltaVault cloud models (AVA-c) are not supported.
Version	4.3
Storage type	Not applicable
Storage amount	Model dependent. Size depends on requirements of your ONTAP Snapshot copy policies to AltaVault

Each AltaVault instance can be connected to one cloud provider. If you configure multiple AltaVault appliances, then each AltaVault appliance can be configured to a different cloud provider. Each AltaVault appliance can support multiple protocols, for example, SnapMirror and SMB.

Each volume in SnapCenter has a relationship with one AltaVault system.

Note: The Data Fabric Solution for Cloud Backup configuration cannot fan out from a single ONTAP volume to multiple AltaVault destinations.

For the latest information and additional AltaVault requirements, see the AltaVault documentation and Interoperability Matrix.

Related information

[NetApp Interoperability Matrix Tool](#)

[NetApp Documentation: AltaVault](#)

Supported browsers

Data Fabric Solution for Cloud Backup supports the following web browsers.

- Chrome version 53 or later
- Internet Explorer 11.0 or later
 - Only default-level security is supported.
Making changes to Internet Explorer security settings results in significant browser display issues.
 - Internet Explorer compatibility view must be disabled.
- Microsoft Edge 2016

Connections to AltaVault and SnapCenter require HTTPS.

For the latest information, see the Interoperability Matrix.

Related information

[NetApp Interoperability Matrix Tool](#)

Licensing requirements

The Data Fabric Solution for Cloud Backup requires several licenses to enable protection operations.

License	Description	Where required
ONTAP	<p>ONTAP requires one of the following licenses:</p> <ul style="list-style-type: none"> • SnapMirror license (by itself or with Premium bundle) • SnapVault license (by itself or with Premium bundle) <p>Some platforms require the Premium bundle; no separate licenses are available for SnapMirror or SnapVault with these platforms.</p> <p>The Premium bundle can be on a legacy system or on a new platform.</p>	

License	Description	Where required
SnapMirror	A license for mirroring backup sets to a destination storage system. Required for this solution. Existing SnapMirror licenses can be used.	On primary storage system only. Typically for FAS-to-FAS replication, SnapMirror is required on both primary and destination systems. With Data Fabric Solution for Cloud Backup, the SnapMirror license is required only on primary storage, not on the destination system because AltaVault is the destination.
SnapVault	An optional license for disk-to-disk backup replication to a destination storage system.	Depends on the type of operation: <ul style="list-style-type: none"> For primary storage (AFF or FAS) to AltaVault, required on primary storage system only. For primary storage (AFF or FAS) to secondary storage (FAS) to AltaVault, required on both primary and destination.
SnapRestore	A required license that enables SnapCenter to restore and verify backup sets.	On primary storage systems. Also required on SnapVault destination systems to perform remote verification and to restore from a backup. Also required on SnapMirror destination systems to perform remote verification.
AltaVault	A capacity-based license required for virtual models to enforce capacity-based storage sizes. Physical appliances do not require a license.	On AltaVault virtual appliance

License	Description	Where required
SnapCenter	<p>SnapCenter includes these licenses:</p> <ul style="list-style-type: none"> Standard – Support for backup and recovery of ONTAP storage, Clone Life Cycle management, basic reporting, task automation, host file systems (Windows, Linux, UNIX), support for custom applications or databases, update of Snapshot copies to SnapMirror and SnapVault secondary destinations, virtualization with VMware, and enterprise applications (Microsoft SQL Server, Oracle) <ul style="list-style-type: none"> FAS and All Flash FAS (AFF): Included in Premium Bundle ONTAP Select and ONTAP Cloud: A la carte license, charged by capacity on the primary storage that is managed by SnapCenter <p>Note: ONTAP Select and ONTAP Cloud are not supported in the Data Fabric Solution for Cloud Backup.</p> Advanced - Optional license. Required to support for NAS shares backup to the cloud using AltaVault and the file catalog <ul style="list-style-type: none"> FAS, AFF, ONTAP Select, ONTAP Cloud: A la carte license for unlimited hosts, charged by capacity on the primary storage managed by SnapCenter <p>Note: A 90 days trial license is available for both SnapCenter standard and advanced licenses.</p> <p>You can use SnapCenter to configure multiple replication relationships with one or more AltaVault appliances without additional licenses.</p> 	On SnapCenter Server
FlexClone	Required for Data Fabric Solution for Cloud Backup.	On primary storage systems
Protocols	For SMB shares, the CIFS license	On primary storage systems. Required on SnapMirror destination systems to serve data if a source volume is unavailable.

Connection and port requirements

Connections and ports must meet minimum requirements before you can install Data Fabric Solution for Cloud Backup.

Data Fabric Solution for Cloud Backup uses the following default ports. You can configure these ports during the installation process.

Type of port	Default port
SnapCenter port	8146 (HTTPS), customizable as in the URL <code>https://server:8146</code>
SnapCenter SMCORE communication port	8145 (HTTPS), bidirectional, customizable
SnapCenter file catalog ports on Linux server	8145 customizable 2181, 2888, 3888, 8983
ONTAP to AltaVault	5010. Requires one or more intercluster LIFs. Recommend connection to 10 Gigabit Ethernet switch.
AltaVault	5010. Requires one or more data interfaces. Only HTTPS is supported for graphical user interface.

Firewalls, proxies, or other network devices should not interfere with connections.

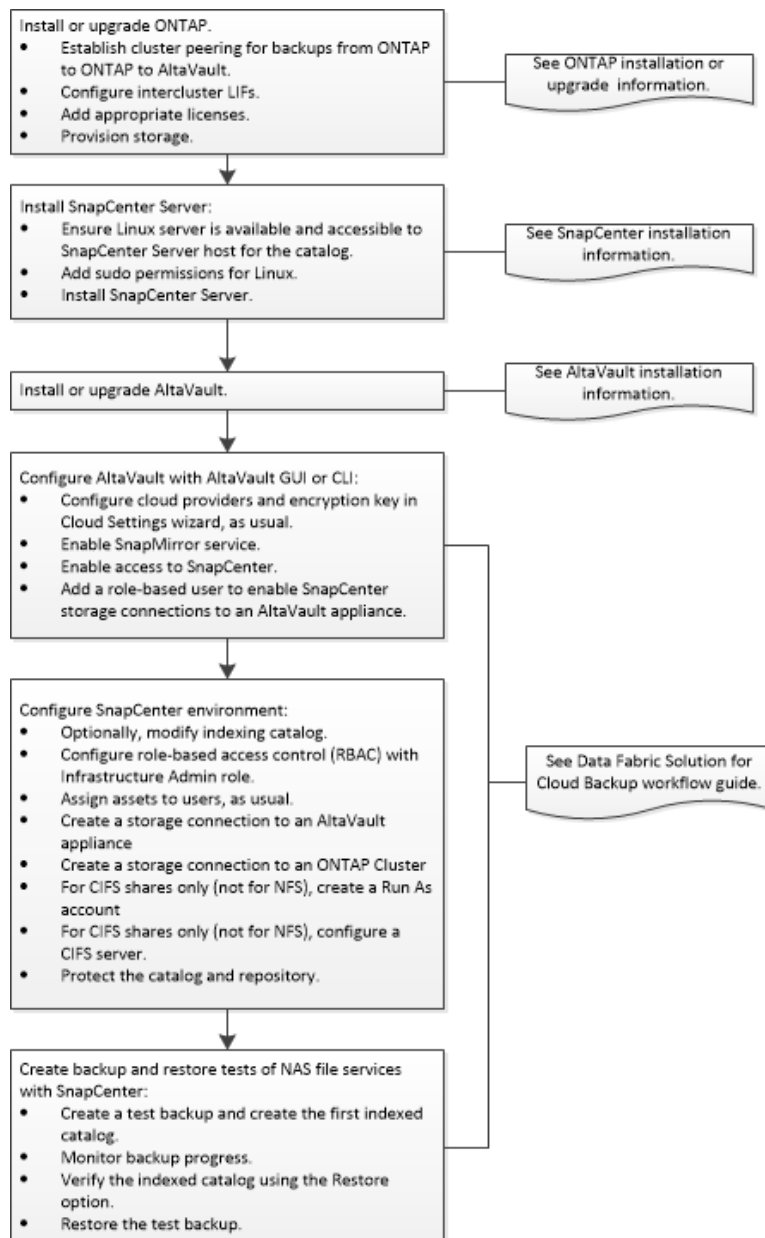
You cannot modify the port number after you install. To change a port number after installation, you must uninstall and install SnapCenter again.

Credential requirements

Using SnapCenter in the Data Fabric Solution for Cloud Backup requires a user with ONTAP cluster administration credentials. These credentials are also required to sign in to the ONTAP CLI, if you prefer to use the CLI.

Installation and configuration workflow with SnapCenter data protection management

The Data Fabric Solution for Cloud Backup can be implemented with or without SnapCenter. Follow this workflow when using SnapCenter.



Related information

[SnapCenter Software 2.0 Installation and Setup Guide](#)

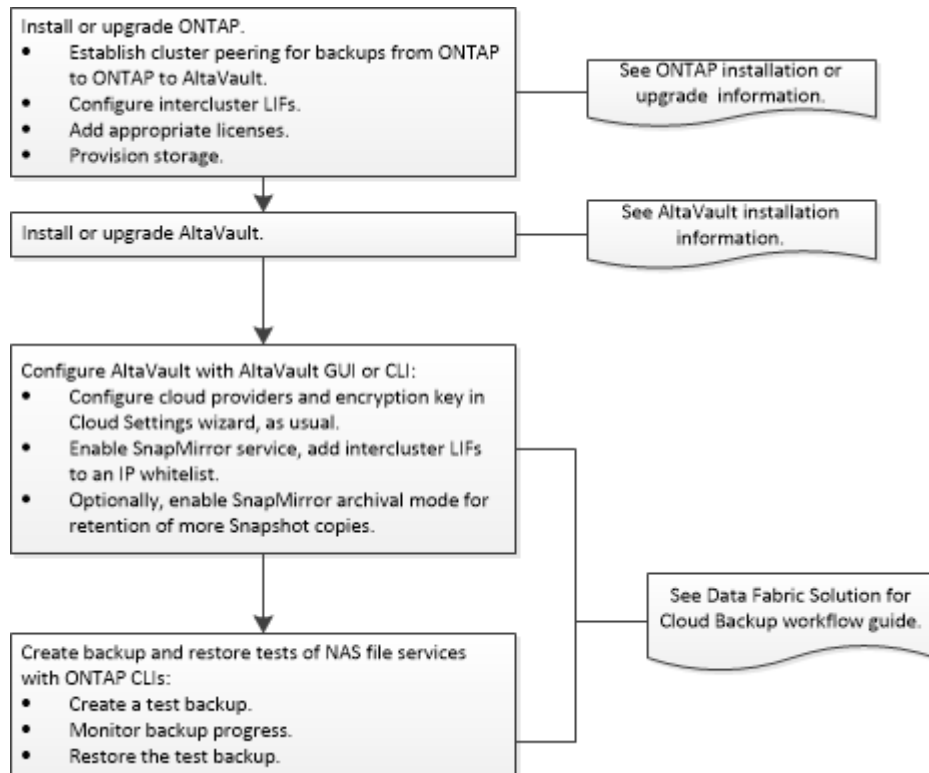
[NetApp Documentation: AltaVault](#)

[ONTAP 9 Cluster Peering Express Guide](#)

[ONTAP 9 Upgrade Express Guide](#)

Installation and configuration workflow without SnapCenter data protection management

The Data Fabric Solution for Cloud Backup can be implemented using ONTAP, AltaVault, and SnapCenter or the solution can be implemented with just ONTAP and AltaVault without SnapCenter data protection management. To install and configure the components of the solution without SnapCenter, you will need to follow the recommended workflow sequence.



Related information

[NetApp Documentation: AltaVault](#)
[ONTAP 9 Cluster Peering Express Guide](#)
[ONTAP 9 Upgrade Express Guide](#)

Installing the Data Fabric Solution for Cloud Backup

Although you can install the components of the solution in a different order, you should follow the recommended installation order.

- ONTAP 9.1. This can be an existing ONTAP deployment or an upgrade to ONTAP 9.1.
- SnapCenter 2.0. This must be a new installation.
- AltaVault 4.3. Typically, you do not use SMB/NFS and SnapMirror deployments on the same AltaVault instance, so we recommend a new AltaVault installation.

Upgrading ONTAP

For the Data Fabric Solution for Cloud Backup, you can use an existing ONTAP deployment of ONTAP 9.1 or upgrade to ONTAP 9.1.

Steps

1. Download the ONTAP upgrade package from the NetApp Support Site at mysupport.netapp.com.
2. Follow ONTAP upgrade instructions.

[ONTAP 9 Upgrade Express Guide](#)

<http://docs.netapp.com/ontap-9/index.jsp>

3. If you intend to back up NAS file services data from ONTAP to AltaVault, add a SnapMirror relationship on the source only. If you intend to back up NAS file services data from ONTAP to ONTAP to AltaVault, add a SnapMirror license on both the source and destination cluster.

- a. Identify licensed features on the source cluster:

```
source_cluster::>system license show
Serial Number: 1-80-000011
Owner: source_cluster
Package      Type      Description      Expiration
-----
Base         site      Cluster Base License  -
NFS          site      NFS License       -
CIFS         site      CIFS License       -
iSCSI        site      iSCSI License      -
FCP          site      FCP License        -
SnapMirror   site      SnapMirror License  -
FlexClone    site      FlexClone License   -
7 entries were displayed.
```

- b. Identify licenses on the destination cluster:

```
destination_cluster::>system license show
```

- c. If any required feature or protocol is not licensed on the source or destination cluster, then add the license.

```
destination_cluster::>system license add -license-code
xxxxxxxxxxxxxx
```

Establishing cluster peering for Data Fabric Solution for Cloud Backup

If you intend to back up NAS file services from ONTAP to ONTAP to AltaVault, Data Fabric Solution for Cloud Backup requires a cluster peering relationship, a one-time operation that connects typically at least three clusters together to enable replication to occur among them.

About this task

You establish the cluster peering relationship for the ONTAP to ONTAP path by completing the following tasks:

- Preparing each cluster
- Creating intercluster logical interfaces (LIFs) on each node
- Setting up each peer relationship

Step

1. Establish the cluster peering relationship by following the cluster peering instructions.

[*ONTAP 9 Cluster Peering Express Guide*](#)

Configuring intercluster LIFs to use dedicated intercluster ports

Configuring intercluster LIFs to use dedicated data ports enables greater bandwidth than using shared data ports on your intercluster networks.

About this task

Creating intercluster LIFs that use dedicated ports involves creating a failover group for the dedicated ports and assigning LIFs to those ports. In this procedure, a two-node cluster exists in which each node has two data ports that you have added, e0e and e0f. These ports are ones you will dedicate for intercluster replication and currently are in the default IPspace. These ports will be grouped together as targets for the intercluster LIFs you are configuring.

Typically, you must configure intercluster LIFs on the peer cluster before you can create cluster peer relationships. However, for the Data Fabric Solution for Cloud Backup, you do not need to create cluster peer relationships after you configure intercluster LIFs.

In your own environment, you might replace the ports, networks, IP addresses, subnet masks, and subnets with those specific to your environment.

Steps

1. List the ports in the cluster by using the `network port show` command.

Example

```
cluster01::> network port show
```

Node	Port	IPspace	Broadcast Domain	Link	MTU	Speed (Mbps) Admin/Oper
cluster01-01						
	e0a	Cluster	Cluster	up	1500	auto/1000
	e0b	Cluster	Cluster	up	1500	auto/1000
	e0c	Default	Default	up	1500	auto/1000
	e0d	Default	Default	up	1500	auto/1000
	e0e	Default	Default	up	1500	auto/1000
	e0f	Default	Default	up	1500	auto/1000
cluster01-02						
	e0a	Cluster	Cluster	up	1500	auto/1000
	e0b	Cluster	Cluster	up	1500	auto/1000

e0c	Default	Default	up	1500	auto/1000
e0d	Default	Default	up	1500	auto/1000
e0e	Default	Default	up	1500	auto/1000
e0f	Default	Default	up	1500	auto/1000

2. Determine whether any of the LIFs are using ports that are dedicated for replication by using the `network interface show` command.

Example

Ports e0e and e0f do not appear in the following output; therefore, they do not have any LIFs located on them:

```
cluster01::> network interface show -fields home-port,curr-port
vserver lif home-port curr-port
-----
Cluster cluster01-01_clus1 e0a e0a
Cluster cluster01-01_clus2 e0b e0b
Cluster cluster01-02_clus1 e0a e0a
Cluster cluster01-02_clus2 e0b e0b
cluster01
cluster01 cluster_mgmt e0c e0c
cluster01
cluster01 cluster01-01_mgmt1 e0c e0c
cluster01
cluster01 cluster01-02_mgmt1 e0c e0c
```

3. If a LIF is using a port that you want dedicated to intercluster connectivity, migrate the LIF to a different port.
 - a. Migrate the LIF to another port by using the `network interface migrate` command.

Example

The following example assumes that the data LIF named `cluster01_data01` uses port e0e and you want only an intercluster LIF to use that port:

```
cluster01::> network interface migrate -vserver cluster01
-lif cluster01_data01 -dest-node cluster01-01 -dest-port e0d
```

- b. You might need to modify the migrated LIF home port to reflect the new port where the LIF should reside by using the `network interface modify` command.

Example

```
cluster01::> network interface modify -vserver cluster01
-lif cluster01_data01 -home-node cluster01-01 -home-port e0d
```

4. Group the ports that you will use for the intercluster LIFs by using the `network interface failover-groups create` command.

Example

```
cluster01::> network interface failover-groups create -vserver cluster01
-failover-group intercluster01 -targets cluster01-01:e0e,cluster01-01:e0f,
cluster01-02:e0e,cluster01-02:e0f
```

5. Display the failover group you created by using the `network interface failover-groups show` command.

Example

```
cluster01::> network interface failover-groups show
Vserver Group Failover
-----
Cluster Targets
```

cluster01	Cluster	cluster01-01:e0a, cluster01-01:e0b, cluster01-02:e0a, cluster01-02:e0b
	Default	cluster01-01:e0c, cluster01-01:e0d, cluster01-02:e0c, cluster01-02:e0d, cluster01-01:e0e, cluster01-01:e0f cluster01-02:e0e, cluster01-02:e0f
	intercluster01	cluster01-01:e0e, cluster01-01:e0f cluster01-02:e0e, cluster01-02:e0f

6. Create an intercluster LIF on the admin SVM cluster01 by using the network interface create command.

Example

This example uses the LIF naming convention adminSVMname_icl# for the intercluster LIF:

```
cluster01::> network interface create -vserver cluster01 -lif cluster01_icl01
               -role intercluster -home-node cluster01-01 -home-port e0e
               -address 192.0.2.250 -netmask 255.255.255.0 -failover-group intercluster01

cluster01::> network interface create -vserver cluster01 -lif cluster01_icl02
               -role intercluster -home-node cluster01-02 -home-port e0e
               -address 192.0.2.251 -netmask 255.255.255.0 -failover-group intercluster01
```

7. Verify that the intercluster LIFs were created properly by using the network interface show command.

Example

```
cluster01::> network interface show
```

Vserver	Logical Interface	Status Admin/Oper	Network Address/Mask	Current Node	Current Port	Is Home
Cluster						
cluster01	cluster01-01_clus_1	up/up	192.0.2.xxx/24	cluster01-01	e0a	true
	cluster01-01_clus_2	up/up	192.0.2.xxx/24	cluster01-01	e0b	true
	cluster01-02_clus_1	up/up	192.0.2.xxx/24	cluster01-01	e0a	true
	cluster01-02_clus_2	up/up	192.0.2.xxx/24	cluster01-01	e0b	true
	cluster_mgmt	up/up	192.0.0.xxx/24	cluster01-01	e0c	true
	cluster01_icl01	up/up	192.0.1.201/24	cluster01-01	e0e	true
	cluster01_icl02	up/up	192.0.1.202/24	cluster01-02	e0e	true
	cluster01-01_mgmt1	up/up	192.0.0.xxx/24	cluster01-01	e0c	true
	cluster01-02_mgmt1	up/up	192.0.0.xxx/24	cluster01-02	e0c	true

8. Verify that the intercluster LIFs are configured for redundancy by using the network interface show command with the -role intercluster and -failover parameters.

Example

The LIFs in this example are assigned the e0e home port on each node. If the e0e port fails, the LIF can fail over to the e0f port.

```
cluster01::> network interface show -role intercluster -failover
```

Vserver	Logical Interface	Home Node:Port	Failover Policy	Failover Group
cluster01-01	cluster01-01_icl01	cluster01-01:e0e	local-only	intercluster01
		Failover Targets:	cluster01-01:e0e, cluster01-01:e0f	
cluster01-01	cluster01-01_icl02	cluster01-02:e0e	local-only	intercluster01
		Failover Targets:	cluster01-02:e0e, cluster01-02:e0f	

9. Display the routes in the cluster by using the `network route show` command to determine whether intercluster routes are available or you must create them.

Creating a route is required only if the intercluster addresses in both clusters are not on the same subnet and a specific route is needed for communication between the clusters.

Example

In this example, no intercluster routes are available:

```
cluster01::> network route show
Vserver  Destination      Gateway      Metric
-----  -
Cluster
cluster01 0.0.0.0/0        192.0.2.1    20
          0.0.0.0/0        192.0.2.1    10
```

10. If communication between intercluster LIFs in different clusters requires routing, create an intercluster route by using the `network route create` command.

The gateway of the new route should be on the same subnet as the intercluster LIF.

Example

In this example, 192.168.1.1 is the gateway address for the 192.168.1.0/24 network. If the destination is specified as 0.0.0.0/0, then it becomes a default route for the intercluster network.

```
cluster01::> network route create -vserver cluster01
-destination 0.0.0.0/0 -gateway 192.0.2.1 -metric 40
```

11. Verify that you created the routes correctly by using the `network route show` command.

Example

```
cluster01::> network route show
Vserver  Destination      Gateway      Metric
-----  -
Cluster
cluster01 0.0.0.0/0        192.0.2.1    20
          0.0.0.0/0        192.0.2.1    10
          0.0.0.0/0        192.0.2.2    40
```

12. If using peer clusters, repeat these steps to configure intercluster networking in the peer cluster.

13. Verify that the ports have access to the proper subnets, VLANs, and so on.

Dedicating ports for replication in one cluster does not require dedicating ports in all clusters; one cluster might use dedicated ports, while the other cluster shares data ports for intercluster replication.

Adding an ONTAP license

A license is a record of one or more software entitlements. Installing license keys, also known as *license codes*, enables you to use certain features or services on your cluster. Data ONTAP enables you to manage feature licenses and monitor feature usage and license entitlement risk.

About this task

Each cluster requires a cluster base license key, which you can install either during or after the cluster setup. Some features require additional licenses.

You can find license keys for your initial or add-on software orders at the NetApp Support Site under **My Support > Software Licenses** (login required). If you cannot locate your license keys from the Software Licenses page, contact your sales or support representative.

Step

1. Add a license:

```
system license add -license-code code
```

Related information

[ONTAP 9 System Administration Reference](#)

Installing SnapCenter for Data Fabric Solution for Cloud Backup

SnapCenter is one of the Data Fabric Solution for Cloud Backup components. Installing SnapCenter involves several tasks.

About this task

Installing SnapCenter for Data Fabric Solution for Cloud Backup involves the following tasks:

- Ensuring that the Linux server is available and accessible to the SnapCenter Server host for the NAS file services catalog
- Adding sudo permissions for Linux
- Installing the SnapCenter Server

Ensuring the Linux server is accessible for NAS file services file catalog

Before you install the SnapCenter Server, you must ensure that the Linux server that you will use for the indexed file catalog is available and will be accessible to the SnapCenter Server.

Steps

1. Before installing the SnapCenter Server, mount the volume for storing the file catalog installation on the following path on the Linux server: `/opt/NetApp/snapcenter/indexer`

This step is optional. It should be completed when you want to allocate dedicated storage for the file catalog, which can take a large amount of space depending on your configuration.

2. Determine whether the Linux server is accessible. Choose one of the following methods:

```
ping <hostname@domainname>
```

```
ping <xxx.xxx.xxx.xxx>
```

3. Use a utility such as PuTTY to determine whether you can connect as the same user that SnapCenter will use to communicate with the catalog.

Configuring sudo privileges for non-root user

SnapCenter 2.0 enables a non-root user to install the Plug-in for Oracle Database and start the plug-in process. However, you must configure sudo privileges for the non-root user to provide access to

several paths. You can configure sudo privileges for the non-root user also if you deploy SnapCenter in the Data Fabric Solution for Cloud Backup, since the file catalog uses a Linux server.

About this task

You must configure sudo privileges for the non-root user to provide access to the following paths:

- /tmp/sc-plugin-installer/snapcenter_linux_host_plugin.bin
- /opt/NetApp/snapcenter/spl/installation/plugins/uninstall
- /opt/NetApp/snapcenter/spl/bin/spl

Steps

1. Create a key by running the following command:


```
dgst -binary -sha224 /tmp/sc-plugin-installer/
snapcenter_linux_host_plugin.bin | openssl base64
```
2. Open the /etc/sudoers file by running the visudo command.
You must use Sudo 1.8.7 or later.
3. Edit the file to include the non-root user name:
 - a. Cmnd_Alias SCCMD = sha224:generated_key /tmp/sc-plugin-installer/
 snapcenter_linux_host_plugin.bin , /opt/NetApp/snapcenter/spl/
 installation/plugins/uninstall, /opt/NetApp/snapcenter/spl/bin/spl
 - b. Defaults:username/groupname !requiretty
 - c. username/groupname ALL=(ALL) NOPASSWD:SETENV: SCCMD

Example

For example, if the non-root user name is 'oracle', modify the lines as:

- Defaults:oracle !requiretty
- oracle ALL=(ALL) NOPASSWD:SETENV: SCCMD

While creating Run As credentials for a non-root user, you must select the **Use sudo privileges** checkbox.

Note: The **Use sudo privileges** checkbox is selected by default for non-root users.

Installing the SnapCenter Server for Data Fabric Solution for Cloud Backup

After you have completed the installation prerequisites, you can use the InstallShield wizard to install the SnapCenter Server.

Before you begin

- Your SnapCenter host system must be up to date with Windows updates with no pending system restarts.
- You must have enabled Windows installer debugging.
See the Microsoft web site for information about enabling and disabling Windows Installer logging.
<https://support.microsoft.com/en-us/kb/223300>
- Your system must meet the prerequisites.

- The SnapCenter Server must be installed on a server that is part of a domain.

Note: SnapCenter Server cannot be installed on a domain controller.

- A two-way trust relationship between the SnapCenter Server domain and the plug-in domain using the Microsoft Active Directory Domains and Trusts snap-in must have been established. While domain trusts, multi-domain forests, and cross-domain trusts are supported, cross-forest domains are not supported.

See Microsoft documentation about Active Directory Domains and Trusts.

<https://technet.microsoft.com/en-us/library/cc770299.aspx>

Steps

1. Optionally, gather information in a SnapCenter installation worksheet provided in the installation instructions.

[SnapCenter Software 2.0 Installation and Setup Guide](#)

2. Download the SnapCenter Server installation package from the NetApp Support Site at mysupport.netapp.com.

3. Install the SnapCenter Server by performing one of the following methods:

- Double-click the downloaded .exe file to launch the SnapCenter Server installer. Proceed through the wizard.
- From a Windows command prompt on the local host, change to the directory where you downloaded the installer and run the .exe file.
If you encounter installation issues, run the installation using a command prompt and generate a log file:

```
SnapCenterversion.exe /debuglog"DirPath\LogFileName"
```

.

For example,

```
SnapCenter2.0.exe /debug"c:\snapcenterlog.txt"
```

The debuglog parameter generates a log file that checks the installation against SnapCenter prerequisites.

4. If you use the wizard, complete the following steps:
 - a. For NAS file services, enter the IP address or host name for one or more servers that you will use to index file catalog information. For high availability, you should add more file catalog servers.
File catalog servers can be configured only during SnapCenter installation.
 - b. For the file catalog server, specify Linux as the server type.
 - c. For the file catalog server, enter port
8145
 - d. For the file catalog server, enter the OS username and password.

Related information

[SnapCenter Software 2.0 Installation and Setup Guide](#)

Installing AltaVault

Data Fabric Solution for Cloud Backup requires AltaVault to perform data protection to the cloud.

About this task

You can use the following deployments of AltaVault:

- Physical
- Virtual (.OVF, Hyper-V, and Linux KVM images)

The AltaVault cloud version (AVA-c) is not currently supported.

Refer to the Interoperability Matrix for the latest information about supported configurations.

Steps

1. Download the AltaVault Physical Appliance or Virtual Appliance package from the NetApp Support Site at mysupport.netapp.com.
2. Do one of the following:
 - For the physical appliance, rack and connect the physical AltaVault appliance according to the AltaVault physical installation guide.
 - For the virtual appliance, deploy the virtual AltaVault appliance according to the AltaVault virtual installation guide.
3. Configure AltaVault, including the following options:
 - Cloud providers
 - AltaVault system settings
 - AltaVault interfaces
 - AltaVault licenses
 - AltaVault encryption key

See the *AltaVault Cloud Integrated Storage Administration Guide* on the NetApp Support Site at mysupport.netapp.com.

Related information

[NetApp Documentation: AltaVault](#)

[NetApp Interoperability Matrix Tool](#)

Configuring AltaVault for Data Fabric Solution for Cloud Backup

Using the Data Fabric Solution for Cloud Backup requires configuring AltaVault. You must have one AltaVault instance per cloud provider with the provider associated with one bucket.

About this task

AltaVault configuration for Data Fabric Solution for Cloud Backup includes the following tasks depending on whether you are using SnapCenter in the deployment:

Tasks	Using SnapCenter in deployment	Not using SnapCenter in deployment
Enable SnapMirror service in AltaVault	Yes	Yes
Add approved intercluster LIFs to an IP whitelist in AltaVault	No	Yes
Optionally, enable SnapMirror long-term retention mode in AltaVault	Yes	Yes
Enable SnapCenter access to AltaVault	Yes	No
Add a role-based user in AltaVault to enable SnapCenter storage connection to an AltaVault appliance	Yes	No
Start the AltaVault Storage Optimization Service	Yes	Yes

Logging in to AltaVault

You can log in to AltaVault using either the graphical user interface (GUI) or a command line interface (CLI).

Steps

- Log in to AltaVault using one of the following:
 - From a web browser, log in to the AltaVault Management Console using HTTPS and the IP address or DNS of the AltaVault Appliance.
 - To access the CLI, log in from an SSH client using the IP address or DNS of the AltaVault Appliance.

For example:

```
ssh admin@172.16.4.4
NetApp AltaVault
admin@172.16.4.4's password:
```

- Enter the administrator credentials.

Enabling the SnapMirror service and whitelist IP connections using the AltaVault GUI

AltaVault supports backup and restore operations for ONTAP FlexVol volumes using the SnapMirror service running on the AltaVault appliance. You must enable the SnapMirror service, which then enables ONTAP to create the SnapMirror relationship in AltaVault and enable backup and restore operations of NAS file services. You must also create a whitelist of approved intercluster LIF IP addresses from which AltaVault accepts connections for data protection.

About this task

When the SnapMirror service is enabled, the shares and Snapshot copies that exist on AltaVault are accessible and can be restored. If the SnapMirror service is disabled, the shares and Snapshot copies that exist on AltaVault are not deleted; however, Snapshot copies are not accessible.

The list of authorized IP addresses must be populated prior to initiating a connection from the source system or the connection will be rejected.

If you use SnapCenter as part of this solution, SnapCenter creates the whitelist of approved IP connections and enables the SnapMirror long-term retention mode when you initiate the backup from SnapCenter. In this case, there is no need to create the IP whitelist and enable SnapMirror long-term retention mode by using AltaVault.

Steps

1. Log in to AltaVault.
2. Select **Configure > SnapMirror**.

The screenshot displays the AltaVault configuration interface for SnapMirror. At the top, there is a navigation bar with links: HOME, CONFIGURE, REPORTS, MAINTENANCE, and HELP. The main title is 'SnapMirror' with a help icon. In the top right corner, there are 'Save' and 'Restart' buttons. The configuration is divided into several sections:

- SnapMirror Service:** Shows the status as 'Enabled' in green. There is a 'Disable' button.
- SnapMirror Shares:** Shows 'Total number of Shares : 0'. There is a 'Remove Selected' button and a table with columns: Name, Peer Path, UUID, and Size. The table is empty with the message 'No shares configured.'
- Whitelist IP:** Shows 'Add Whitelist IP' and 'Remove Selected' buttons. There is a table with columns: IP Address, and two IP addresses are listed: 192.168.2.1 and 192.168.1.2.
- Long Term Retention:** Shows the status as 'Disabled' in red. There is an 'Enable' button.
- SnapCenter Access:** Shows the status as 'Disabled' in red. There is an 'Enable' button.

3. Under SnapMirror Service, click **Enable**.
4. If the "Service restart required" prompt appears, click the **Restart** button that becomes enabled in the upper right.

5. Under Whitelist IP, click **Add Whitelist IP**.
6. Enter the IP addresses of the intercluster LIFs from which AltaVault will accept connections for backup and restore operations and click **Add**.

To remove an IP address, select the IP address and click **Remove Selected**.

Note: Removing an IP address from the whitelist disables access to the AltaVault from that IP address.

7. If you are using SnapCenter to initiate and manage data protection operations, under SnapCenter Access, click **Enable**.

This enables communication between AltaVault and SnapCenter.

8. If the Service restart required prompt appears, click **Restart** in the upper right of the console.

Enabling SnapMirror long-term retention mode in AltaVault

AltaVault supports up to 500 shares in one of two modes: short-term retention mode (default) or long-term retention mode. For Data Fabric Solution for Cloud Backup, the mode impacts the retention of your Snapshot copies. SnapMirror long-term retention mode is required if you want to retain more than 251 Snapshot copies for a relationship on AltaVault. Long-term retention mode is disabled by default.

About this task

Snapshot copy retention is dependent upon the retention policy set up in either ONTAP or SnapCenter. For example, if the retention policy is set to 100 Snapshot copies, and the maximum is reached, AltaVault begins deleting the oldest Snapshot copies to make room for newer ones.

- In short-term retention mode, each share supports up to 251 Snapshot copies. The Snapshot copy retention is governed by the SnapMirror policy set in either ONTAP or SnapCenter.
For example, suppose a share has a two-tier retention policy supporting 50 hourly and 100 daily Snapshot copies. In this case, when the count of hourly Snapshots exceeds 50 or the daily count exceeds 100, the oldest Snapshot copy of the respective tier is deleted.
- In long-term retention mode, each share supports up to 3700 Snapshot copies, which is the equivalent to 10 years of daily Snapshot copies. In this mode, AltaVault continues to store Snapshot copies until it reaches the maximum. If a share exceeds 3700 Snapshot copies, AltaVault begins deleting the oldest Snapshot copies to make room for newer ones.

Retention of Snapshot copies in archival mode can be managed by AltaVault or SnapCenter:

- Snapshot copies managed by AltaVault: AltaVault continues accumulating Snapshot copies with a maximum of 3700 before deleting the oldest Snapshot copy in the share. The oldest is deleted irrespective of the label (hourly/daily/monthly).
- Snapshot copies managed by SnapCenter: SnapCenter manages any number of Snapshot copies with a maximum of 3700.

Long-term retention mode is enabled automatically when SnapCenter transfers data to AltaVault.

When long-term retention mode is disabled, AltaVault reverts to using the retention policy set up in ONTAP or SnapCenter, which supports a maximum of 251 Snapshot copies per share. If there are large numbers of Snapshot copies (more than 251) when long-term retention mode is disabled, the number of Snapshot copies will be reduced to match the count set in the retention policy.

Warning: When you use SnapCenter to manage backups, you should not disable long-term retention mode in AltaVault.

The mode used for long-term retention applies to all SnapMirror shares created on AltaVault.

If you use SnapCenter as part of this solution, SnapCenter creates the whitelist of approved IP connections and enables the SnapMirror long-term retention mode when you initiate the backup from SnapCenter. In this case, there is no need to create the IP whitelist and enable SnapMirror long-term retention mode by using AltaVault.

Steps

1. Log in to AltaVault.
2. Under SnapMirror Long-term Retention Mode, click **Enable**.

Enabling SnapCenter access to AltaVault

AltaVault supports using SnapCenter to back up, delete, and restore Snapshot copies, including individual file restores. If you are using SnapCenter to manage backups to AltaVault, you must enable SnapCenter access on AltaVault.

Steps

1. Log in to AltaVault.
2. Select **Configure > SnapMirror**.
3. If you are using SnapCenter to initiate and manage data protection operations, under SnapCenter Access, click **Enable**.

Adding a SnapCenter role-based user account to enable SnapCenter storage connections to an AltaVault appliance

To perform data protection operations using the Data Fabric Solution for Cloud Backup, you must add a SnapCenter role-based account to AltaVault. You will later add a storage connection in SnapCenter and will need this account's credential information.

About this task

You can add a user using the AltaVault graphical user interface (GUI) or by using the AltaVault CLI.

Steps

1. Log in to AltaVault.
2. Select **Configure > User Permissions**.

User Permissions

Capability-Based Accounts:

User	Enabled
admin	true
monitor	false

Role-Based Accounts:

☒ Add a New User ☐ Remove Selected Users

Account Name:

Password:

New Password Confirm:

☐ External Authentication Only

☒ Enable Account

Roles and Permissions:

Setting	Select All	Read-Only	Read/Write
Admin Settings:	<input type="radio"/> Deny	<input type="radio"/> Read-Only	<input checked="" type="radio"/> Read/Write
General Settings:	<input type="radio"/> Deny	<input type="radio"/> Read-Only	<input checked="" type="radio"/> Read/Write
Prepop Settings:	<input type="radio"/> Deny	<input type="radio"/> Read-Only	<input checked="" type="radio"/> Read/Write
Replication Settings:	<input type="radio"/> Deny	<input type="radio"/> Read-Only	<input checked="" type="radio"/> Read/Write
Reports Settings:	<input type="radio"/> Deny	<input type="radio"/> Read-Only	<input checked="" type="radio"/> Read/Write
Security Settings:	<input type="radio"/> Deny	<input type="radio"/> Read-Only	<input checked="" type="radio"/> Read/Write
Storage Settings:	<input type="radio"/> Deny	<input type="radio"/> Read-Only	<input checked="" type="radio"/> Read/Write

3. In the Role-Based Accounts section, click **Add a New User** to add a user account that will be used by SnapCenter.
4. Enter an account name and password and check **Enable Account**.
5. Click **Read/Write** permissions for the following roles: General Settings, Replication Settings, and Storage Settings.
6. Click **Add**.

Starting the AltaVault Storage Optimization Service

After you configure other AltaVault settings, you must start the Storage Optimization Service, which provides compression, deduplication, and encryption of data upon ingest into AltaVault.

Steps

1. From the AltaVault menu, select **Maintenance > Service**.
2. To verify that the service is running again, select **Home > Optimization Service**.

The display indicates that the service is running and that the status is "Ready."

Configuring SnapMirror and enabling SnapCenter access using the AltaVault CLI

In addition to using a graphical user interface, you can use the AltaVault command-line interface (CLI) for configuring SnapMirror. SnapMirror configuration includes enabling the SnapMirror service, restarting the Storage Optimization Service, configuring a whitelist of approved IP addresses from which AltaVault accepts SnapMirror connections, enabling SnapMirror long-term retention mode, and enabling SnapCenter access.

About this task

If you already configured SnapMirror using the AltaVault graphical user interface, you do not need to perform this procedure.

If you use SnapCenter as part of this solution, SnapCenter creates the whitelist of approved IP connections and enables the SnapMirror long-term retention mode when you initiate the backup from

SnapCenter. In this case, there is no need to create the IP whitelist and enable SnapMirror long-term retention mode by using AltaVault.

Steps

1. Use SSH to log in to the AltaVault CLI and enter configuration mode:

```
<hostname> > enable
<hostname> # configuration terminal
<hostname> (config) #
```

2. Enable SnapMirror service:

```
<hostname> (config)# snapmirror enable
```

To disable the SnapMirror service, use the `no snapmirror enable` command.

3. At the prompt, restart the Storage Optimization Service:

```
<hostname> (config)# service restart
```

4. Verify that the SnapMirror service is ready:

```
<hostname> (config)# show snapmirror state
SnapMirror Server: Ready
```

5. If you are not using SnapCenter to manage NAS file services data protection, configure the list of IP addresses from which AltaVault accepts SnapMirror data:

You can enter only one IP address at a time.

```
<hostname> (config)# snapmirror whitelist add ip <IP address>
```

6. Display a list of IP addresses:

```
show snapmirror whitelist
  Entry: 192.0.2.253
  Entry: 192.0.2.254
```

To delete an IP address from the list, enter the `snapmirror whitelist delete ip <IP address>` command.

7. Optionally, display and enable long-term retention mode to increase the number of Snapshot copies that can be retained in AltaVault:

```
<hostname> (config)# show snapmirror archival
SnapMirror Archival mode: disabled
<hostname> (config)# snapmirror archival enable
```

To disable long-term retention mode, enter the `no snapmirror archival` command.

8. Enable SnapCenter access to AltaVault:

```
<hostname> (config)# rest enable
```

To disable SnapCenter access, enter the `no rest enable` command.

9. Verify that SnapCenter access is enabled in AltaVault:

```
<hostname> (config)# show rest status
REST Service: running
```

The status is either running (enabled) or stopped (disabled).

10. Create a role-based user account for SnapCenter in AltaVault:**a. Create the account:**

```
<hostname> (config)# username <username> password <password>
```

b. Set the role-based permissions on the account:

```
<hostname> (config)# rbm user <username>
    role cb_general_settings permission read-write
<hostname> (config # rbm user <username>
    role cb_replication_settings permission read-write
<hostname> (config # rbm user <username>
    role cb_storage_settings permission read-write
```

c. Verify the account:

```
<hostname> (config)# show rbm users
User: storadmin
  role: cb_general_settings      permissions: read-write
  role: cb_replication_settings  permissions: read-write
  role: cb_storage_settings      permissions: read-write
```

11. Restart the Storage Optimization Service:

```
<hostname> (config)# service restart
```

12. Verify that the Storage Optimization Service is running:

```
<hostname> (config)# show service
Storage Optimization Service: ready
```

Configuring SnapCenter for Data Fabric Solution for Cloud Backup

Using the Data Fabric Solution for Cloud Backup requires configuring SnapCenter.

About this task

To configure SnapCenter for Data Fabric Solution for Cloud Backup, you should complete the following:

- Configure the predefined SnapCenter Infrastructure Admin role for AltaVault access and setting up role-based access control
- Configure a storage connection to an AltaVault appliance
- Configure a storage connection to an ONTAP cluster
- For CIFS shares only (not for NFS), create a Run As account
- For CIFS shares only (not for NFS), configure CIFS server credentials
- Protect the SnapCenter file catalog and SnapCenter repository

Logging in to SnapCenter

Through SnapCenter role-based access control, users or groups are assigned roles and resources. When you log in to the SnapCenter graphical user interface, you log in with an Active Directory account.

About this task

The SnapCenter graphical user interface (GUI) URL is configured based on information that you provide during installation. It is useful to know where to find it after you complete the SnapCenter installation.

During the installation, the SnapCenter Server Install wizard creates a shortcut and places it on your local host desktop. Additionally, at the end of the installation, the Install wizard provides the SnapCenter URL, which you can copy if you want to log in from a remote system.

Attention: Closing just the SnapCenter browser tab does not log you off of SnapCenter if you have multiple tabs open in your web browser. If you need to comply with security requirements, you must log off of SnapCenter either by clicking the **Sign out** button or shutting down the entire web browser.

Attention: Do not allow your browser to save your SnapCenter password. Saving your password creates a high security risk.

The default GUI URL is a secure connection to port 8146 on the server where the SnapCenter Server is installed (`https://server:8146`). If you provided a different server port during the SnapCenter installation, that port is used instead.

For Network Load Balance (NLB) deployment, you must access SnapCenter using the NLB cluster IP (`https://NLB_Cluster_IP:8146`).

In addition to using the SnapCenter GUI, you can use the following interfaces:

- PowerShell cmdlets for Microsoft SQL databases and Windows file systems

You can use PowerShell cmdlets on SQL databases to script configuration, backup, restore, verification, and clone operations. You can use PowerShell cmdlets on Windows file systems to script backup, restore, and clone operations.

- SnapCenter command-line interface (CLI), *sccli* for Oracle databases on Linux machines. You can use the CLI to script configuration, backup, catalog, uncatalog, restore, verification, mount, unmount, and clone operations.

Note: Some cmdlets have changed in SnapCenter 2.0. If you use cmdlets in scripts, you might need to update your scripts.

For details, see the SnapCenter cmdlet or SnapCenter CLI documentation.

Steps

1. Launch SnapCenter from the shortcut located on your local host desktop, from the URL provided at the end of the installation, or from the URL provided to you by your SnapCenter administrator.
2. Choose which user you want to log in as:

If you want to ...	Do the following ...
Log in as the SnapCenter administrator	Enter the domain user with local administrator credentials provided during the SnapCenter installation.
Log in as a SnapCenter user	Enter your user credentials: <i>Domain\UserName</i>

3. If you are assigned more than one role, from the **Role** box, select the role you want to use for this login session.

Your current user and associated role are shown in the upper right of SnapCenter.

Configuring the Infrastructure Admin role for AltaVault access

To configure role-based access control for SnapCenter users, you complete a two-step process: you add users or groups to a role and then add assets to that user. The first step of adding users or groups to a role determines the options that SnapCenter users can access.

Before you begin

You must have logged in as the SnapCenterAdmin role.

About this task

The Data Fabric Solution for Cloud Backup requires that you add users to the predefined Infrastructure Admin role or create a role using the same permissions. When you assign users to this role, you should assign the AltaVault administration account for authorization and authentication. The Infrastructure Admin role enables users assigned to this role access to the SnapCenter file catalog.

Steps

1. In the left navigation pane, click **Settings**.
2. In the **Settings** page, click **Roles**.
3. In the **Roles** page, select the Infrastructure Admin role to which you want to add the user.

4. Click **Modify**.
5. Click **Next** until you reach the **Users/Groups** page of the wizard.
6. In the **Users/Groups** page, specify the domain to which the user belongs.
7. In the **user or group name** field, enter a user or group name and click **Add**.
Repeat to add additional users or groups to the selected role.
8. Click **Next** to view the summary, and then click **Finish**.

After you finish

To continue with role-based access control, assign assets to users next.

Assigning users access to assets

Setting up role-based access control (RBAC) for users is a two-step process. After you add a user or group to a role that contains the appropriate permissions, you must assign the user access to SnapCenter assets, such as hosts and storage connections. This enables users to perform the actions for which they have permissions on the assets that are assigned to them.

About this task

You can assign access to users even if the user is not part of a role. This helps you add users; however, you must add the user to a role at some point to take advantage of role-based access control efficiencies.

You can also assign Run As credential maintenance to a user.

If you are planning to replicate Snapshot copies to a mirror or vault, you must assign the storage connection for both the source and destination volume to the user performing the operation.

You should add assets before assigning access to the users.

Steps

1. In the left navigation pane, click **Settings**.
2. In the **Settings** page, click **SnapCenter Assets**.
3. In the **My SnapCenter Assets** page, select the type of asset you want to assign from the **SnapCenter Asset Type** drop-down list.
4. In the **SnapCenter Asset Name** table, highlight the asset you want to assign and click **Assign**.
5. Provide the domain name to which the asset belongs.
6. Choose whether you want to assign access to a user or a group.
7. Enter the user or group name and click **Add**.
If you are not sure about the name, use a more advanced search by clicking the **Search users or group** heading, typing a partial name, and clicking **Search**.
8. Repeat this procedure until each user or group has all the required assets.
9. Click **OK** to save your changes.

Setting up storage system connections

Before you can perform backup, restore, clone, and provisioning operations with SnapCenter, you must set up the storage system connections that give SnapCenter access to ONTAP storage. If you are configuring connections for the Data Fabric Solution for Cloud Backup, you must create connections to each AltaVault system and one to an ONTAP Cluster.

Before you begin

You must have permissions in the Infrastructure Admin role to create storage connections.

About this task

If you are planning to replicate Snapshot copies to a SnapMirror or SnapVault destination, make sure to set up storage system connections for the destination volume as well as the source volume.

Steps

1. In the left navigation pane, click **Storage Systems**.
2. In the **Storage Systems** page, click **New**.

New Storage Connection ×

Provide information for the Storage System you want to add.

Storage System

Name or IP

User name

Storage System user name

Password

Storage System password

Storage Type

ONTAP Cluster ▾

Site

Site

Protocol

HTTPS ▾

Port

443

Timeout

60

seconds ⓘ

☐ Preferred IP address ⓘ

Event Management System (EMS) & AutoSupport Settings

☐ Send EMS messages to log SnapCenter Server events to storage system syslog

☐ Send AutoSupport notification for failed operations to storage system

Cancel

OK

3. In the **New Storage Connection** wizard, provide the following information:

For this field...	Do this...
Storage System	<p>Enter the storage system name or IP address.</p> <p>Note: Storage system names, not including the domain name, must be 15 characters or fewer. To create storage system connections with names with more than 15 characters, you can use the <code>Add-SmStorageConnection</code> PowerShell cmdlet.</p> <p>SnapCenter does not support multiple SVMs with the same name on different clusters. Each SVM supported by SnapCenter must have a unique name.</p>
User name/Password	<p>Do one of the following:</p> <ul style="list-style-type: none"> • ONTAP: Enter the credentials used (usually VSAdmin) to access the storage system. • AltaVault: Enter the role-based access control credentials you entered in AltaVault.
Storage Type	<p>Select ONTAP SVM, ONTAP Cluster, or AltaVault.</p> <p>If you are configuring connections for SnapCenter application plug-ins, choose ONTAP SVM as the storage type.</p> <p>If you are configuring connections for the Data Fabric Solution for Cloud Backup, you must add connections to each AltaVault system and to an ONTAP Cluster. For each AltaVault system, choose AltaVault as the storage type, and choose ONTAP Cluster as the second storage type.</p> <p>Note: If you add a new ONTAP SVM connection and you have already added an ONTAP Cluster for the Data Fabric Solution for Cloud Backup, after you add the SVM, you must run the Modify Storage Connection wizard on the ONTAP Cluster that you already added to ensure the following updates are made:</p> <ul style="list-style-type: none"> • The SVM is in the NSM database • CIFS server discovery is enabled on the SVM • The SnapCenter cache is updated with information about the new SVM
Site	<p>Applicable only if you choose ONTAP Cluster or AltaVault as your storage type. Enter the physical site name, for example, the data center city. The site you enter is displayed in the SnapCenter interface.</p>
Protocol	<p>Select the protocol used for connection to the SVM that was configured during SVM setup, typically HTTPS.</p>
Port	<p>Enter the port that the storage system accepts.</p> <p>The defaults typically work.</p> <p>If you are configuring the AltaVault connection, enter the AltaVault port.</p>

For this field...	Do this...
Timeout	Applicable only if you choose ONTAP Cluster or ONTAP SVM as the storage type. Enter the time in seconds that should elapse before communication attempts are halted. The default value is 60 seconds.

If you have questions about these values, consult your storage administrator.

- Optional: If the SVM has multiple management interfaces, select the **Preferred IP address** check box, and then enter the preferred IP address for SVM connections.

Note: If you have more than one iSCSI or FC session configured per SVM connected to the host, then use the host device for multipathing rather than using the native device. While performing a restore, clone, mount, or backup verification operation, if the storage system and the host have iSCSI and FC configured together, then FC is preferred.

ONTAP 9 Cluster Management Using OnCommand System Manager

- Applicable only if you choose ONTAP Cluster or ONTAP SVM as the storage type. If you want to have Event Management System (EMS) messages sent to the storage system syslog or have AutoSupport messages sent to the storage system for failed operations, select the appropriate check box.

When you select the AutoSupport check box, the EMS messages check box is also selected because EMS messaging is required to enable AutoSupport notifications.

- Click **OK**.

Setting up your Run As credentials

SnapCenter uses Run As credentials to authenticate users for SnapCenter operations. You should create Run As credentials for installing SnapCenter plug-ins and additional Run As credentials for performing data protection operations on databases or Windows file systems.

About this task

- Linux hosts
You must set up Run As credentials for scheduling on Linux hosts.
Although you are allowed to create Run As credentials for Linux after deploying hosts and installing plug-ins, the best practice is to create Run As credentials after you add SVMs, before you deploy hosts and install plug-ins. You must set up the Run As credentials for the root user or a non root user who has sudo privileges to install and start the plug-in process.
- Windows hosts
You must set up Windows Run As credentials before installing plug-ins.
Set up the Run As credentials with administrator privileges, including administrator rights on the remote host.
- SQL authentication on Windows hosts
You must set up SQL Run As credentials after installing plug-ins.
If you are deploying SnapCenter Plug-in for Microsoft SQL Server, you must set up SQL Run As credentials after installing plug-ins. Set up a Run As credential for a user with SQL Server sysadmin permissions.
The SQL authentication method authenticates against a SQL Server instance. This means that a SQL Server instance must be discovered in SnapCenter. Therefore, before adding a SQL Run As credential, you must add a host, install plug-in packages, and refresh resources. You need SQL Server authentication for performing operations such as scheduling or discovering resources.

- Custom Plug-ins applications

The plug-in uses the Run As credentials selected or created when adding a resource. If a resource does not require a Run As credentials during data protection operations, you can set the Run As credentials as **None**.

- For Data Fabric Solution for Cloud Backup for CIFS shares only (not for NFS), you must create Run As credentials and configure a CIFS server for the ONTAP cluster.

Steps

1. In the left navigation pane, click **Settings**.
2. In the **Settings** page, click **Run As Credentials**.
3. Click **New**.

Run As Credentials [X]

Provide information for the Run As Credentials you want to add

Run As name: NonRootUser

Username: oracle [i]

Password: *****

Authentication: Linux [v]

☒ Use sudo privileges [i]

Cancel OK

4. In the **Run As Credentials** page, do the following:

For this field...	Do this...
Run As name	Enter a name for the Run As credentials.
User name/Password	Enter the user name and password used for authentication. You must add the domain name as the prefix to the username.

For this field...	Do this...
Authentication Mode	<p>Select the authentication mode that you want to use.</p> <p>If you select the SQL authentication mode, you must also specify the SQL server instance and the host where the SQL instance is located.</p> <p>For Data Fabric Solution for Cloud Backup for CIFS shares only, select Windows as the mode.</p>
Use sudo privileges	Applicable to Linux users only. Select the Use sudo privileges check box if you are creating run as credentials for a non-root user.

5. Click **OK**.

After you finish

After you finish setting up Run As credentials, you might want to assign Run As credential maintenance to a user or group of users on the My SnapCenter Assets page.

Configuring CIFS servers for Data Fabric Solution for Cloud Backup

For Data Fabric Solution for Cloud Backup for CIFS shares only (not for NFS), you must create and configure a CIFS server for the ONTAP cluster. This configuration associates the Run As credentials with the CIFS server for authentication.

Before you begin

You must have already fulfilled these prerequisites:

- You must have already created Windows Run As credentials for CIFS.
- The CIFS server domain must already be in a trust relationship with the SnapCenter Server domain.
- The CIFS server name should be available from the SnapCenter Server host and file catalog host.

Steps

1. In the left navigation pane, click **Storage Systems**.
2. In the **Storage Systems** page, click on an ONTAP cluster.
3. Click **Configure CIFS Servers**.
4. For each CIFS server, select the associated Run As credential.
5. Click **OK**.

Protecting the SnapCenter file catalog and the repository

For the Data Fabric Solution for Cloud Backup solution, SnapCenter includes a NAS file services file catalog, which you should protect. You should also protect the SnapCenter repository. You can

protect both using a PowerShell `protect-SmRepository` cmdlet. If you have not done so previously for SnapCenter use, you should prepare the PowerShell environment.

About this task

During the SnapCenter installation, the SnapCenter repository is placed on the C drive of the Windows host. However, you might want to move it to a NetApp volume. To do so, you can use the SnapCenter `Protect-SmRepository` cmdlet.

Preparing the PowerShell environment

Preparing the PowerShell environment includes verifying that the modules that contain the cmdlets are loaded, and if not, importing the necessary modules.

About this task

For information about PowerShell cmdlets, use the SnapCenter cmdlet help or see the cmdlet reference information.

[SnapCenter Software 2.0 Windows Cmdlet Reference Guide](#)

Steps

1. On either the SnapCenter server or the application host, open a PowerShell session window and verify that the proper modules are loaded by using the `get-module` cmdlet.

If you execute the cmdlet on the SnapCenter server, only the SnapCenter module is displayed. The cmdlets that have “Sd” in the suffix, which are the cmdlets that reside in the SnapDrive module, are available only on the application host.

Example

```
PS C:\> get-module
```

ModuleType	Version	Name	ExportedCommands
Manifest	3.1.0.0	Microsoft.PowerShell.Management	{Add-Computer, Add-Content, Checkpoint-Computer, Clear-Con...}
Manifest	3.1.0.0	Microsoft.PowerShell.Utility	{Add-Member, Add-Type, Clear-Variable, Compare-Object...}
Manifest	1.0	SnapCenter	{Add-SmResourceGroup, Add-SmGroupToRol...}
Manifest	1.0.0	SnapDrive	{Add-SdIgroupInitiator, Add-SdLunMap, Add-SdPortSetPort, C...}

2. If the cmdlets that reside on the SnapDrive module are not displayed, import that module:
`import-module`

Note: `import-module` imports a module only into the current session. To import a module into all sessions, add an `import-module` cmdlet to your Windows PowerShell profile.

Example

```
PS C:\> get-module -listavailable snap* | import-module
PS C:\> get-module
```

Directory: C:\Windows\system32\WindowsPowerShell\v1.0\Modules

ModuleType	Version	Name	ExportedCommands
-----	-----	----	-----
Manifest	1.0	SnapCenter	{Add-SmPolicy, Add-SmRunAs, Add-SmVeri...
Manifest	1.0.0	SnapDrive	{Add-SdIgroupInitiator, Add-SdPortSetPort, Connect-SdIscsi...

Protecting the file catalog and SnapCenter repository with PowerShell cmdlets

For the Data Fabric Solution for Cloud Backup solution, SnapCenter includes a file catalog, which you should protect. You should also protect the SnapCenter metadata repository. You can accomplish both using a PowerShell cmdlet.

About this task**Steps**

1. Launch PowerShell.
2. From the SnapCenter Server command prompt, enter the following and enter your credentials:

```
Open-SMConnection
```

3. Protect the file catalog by using the `Protect-SmRepository` command:

```
Protect-SmRepository [-HostName] string [-Path] Hashtable [-  
InstanceCredential] [[-SQLServerAuthenticationMode] SQLServerAuthMode] [[-  
RetentionCount] Int32] CommonParameters [-CatalogBackupPath] string  
  
Protect-SmRepository [-HostName] string [-Path] Hashtable [[-  
RetentionCount] Int32] CommonParameters [-CatalogBackupPath] string
```

HostName

Specifies the SnapCenter database host name. If the SnapCenter database is hosted by a failover cluster instance (FCI), then specify the FCI owner host name.

Path

Specifies the NetApp destination disk path.

Schedule

Specifies the backup schedule type.

Options are as follows:

ScheduleType

Specifies the backup schedule type.

StartTime

Specifies the scheduled backup start time. The default is the current time.

EndTime

Specifies the scheduled end time.

RetentionCount

Specifies the number of backups to retain. By default, seven backups are retained.

CatalogBackupPath

Specifies an NFS mount point on ONTAP storage. The NFS mount point must be accessible to all backup catalog servers.

This parameter is used only when you have NAS file services installed.

Example

```
Protect-SmRepository -HostName Dan-NGVM2.sme71lad.net -Path F:\
-Schedule @"{"ScheduleType"="hourly";"StartTime"="11/30/2016 5:55 AM";
"RepeatTask_Every_Hour"="00:15"}" -CatalogBackupPath /mnt/catalog_backup
-RetentionCount 4
```

4. Protect the SnapCenter metadata repository:

```
Protect-SmRepository [-HostName] string [-Path] Hashtable [[-
RetentionCount] Int32] CommonParameters [-CatalogBackupPath] string
```

Example

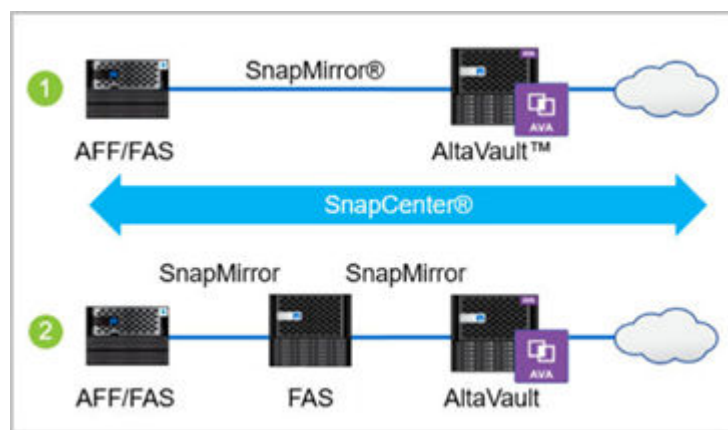
```
Protect-SmRepository -HostName NB-MVA-
DEV057.nbsdsm.mycompany.netapp.in
-Path E:\DBs -InstanceCredential sa -SQLServerAuthenticationMode SQL
-Schedule @"{"ScheduleType"="hourly";"StartTime"="10/21/2016 5:18 PM"}
```

Backing up NAS file services to the cloud and restoring them with SnapCenter management

You can back up NAS file services to the cloud and restore them from the cloud using SnapCenter and AltaVault. You can perform a baseline backup from SnapCenter. Subsequently, all backups of that volume from SnapCenter are incremental backups and are initiated by the policy-driven schedule.

About this task

You can back up all shares on a FlexVol volume from primary or secondary storage to AltaVault, which then replicates data to the cloud, as specified in AltaVault and SnapCenter configurations.



Each volume is linked with one share in AltaVault.

You can define policies that govern the timing and retention of backups.

Later, if needed, you can perform a single file restore of data from AltaVault to ONTAP using SnapCenter.

Note: With the ONTAP CLI, you can perform a full volume restore or a single file restore.

As a best practice, we recommend that you initiate no more than 10 concurrent restore operations. Otherwise, transfer time outs and slower restore throughput may occur. SnapCenter supports up to 100 concurrent transfers to a single AltaVault node, including both backup and restore operations. This means that if 100 concurrent backups are running against a single node, you cannot initiate any restore operations from that node.

You can add up to 80 volumes in a single policy-governed backup operation.

Using policies to govern backing up NAS file services

A backup policy is a set of rules that governs how you manage and retain backups, and how frequently the resource or resource group is backed up. Additionally, you can specify replication and script settings. Specifying options in a policy saves time when you want to reuse the policy for

another resource group. To use a predefined policies, you must first copy and modify it to meet your needs before you can use it to govern data protection.

About this task

SnapCenter includes the following predefined policies that govern the backup of NAS file services:

- Backup to the Cloud: Backs up ONTAP FlexVol volumes on primary storage to AltaVault.
- Backup to Vault to Cloud: Backs up ONTAP FlexVol volumes on primary storage to secondary storage and then to AltaVault.

To back up NAS file services, you must first copy and modify predefined policies to suit your needs.

Note: You cannot use the predefined policies for NAS file services without first copying them and you cannot create your own policies for NAS file services.

You can copy and modify policies before or during the setup of the backup process. This procedure describes how to copy a policy before setting up the backup process.

Steps

1. In the SnapCenter left navigation pane, click **Settings**.
2. On the Settings page, click **Policies**.
3. To review the details of any predefined policies, click the predefined policy and click **Details**.
4. Click the policy, click **Copy**, and enter a new policy name.
5. Click the copy and click **Modify**.
6. In the **Policy** wizard, enter the new policy name, which must be unique, and an optional description and click **Next**.

Backup to Vault to Cloud- Copy2

Configure Protection Policy

1 Name

2 Protection

3 Summary

Local copies: Hourly 24, Daily 7

Vault copies: Daily 30, Weekly 52, Monthly 24

Cloud copies: Daily 2000, Weekly 104, Monthly 84

Vault Site: SVL 121

Choose Storage Cluster: dsg-3250-05-06.sdworld.lab.netapp.com

Choose Storage VM: tan_peer2

Backup Transfer Start Time: 12:00 AM

Transfer Window: 3 hours

Lag Time: 2 hours

Note: Backup creation time values are with respect to server time zone

Previous Next

7. In the **Protection** page, click on one of the protection boxes at the top to configure the policy.

Option	Description
Site or Vault Site	Enter the data center site name.
Cloud Bucket	Applicable only if you select the cloud copy. Enter the name of the bucket associated with the cloud provider.

Option	Description
Choose Storage Cluster	Applicable only if you select the vault copy. Select the ONTAP Cluster as the backup destination.
Choose Storage VM	Applicable only if you select the vault copy. Select the Storage Virtual Machine (SVM) as the backup destination.
Backup Transfer Start Time	Enter the time when the initial baseline backup should begin. This is based on the ONTAP source storage time zone.
Transfer Window	Select the number of hours during which the backup transfer should occur. After that time plus the lag time elapses, if the backup is not available in ONTAP SnapVault or AltaVault, the backup will be in a warning state.
Lag Time	Select the number of additional hours that the backup transfer can take after the transfer window elapses. The minimum lag time is 1 hour. After the transfer window plus the lag time elapses and if the backup is not available in ONTAP SnapVault or AltaVault, the backup will be in a warning state.
Daily, Weekly, Monthly options	Enter the time when the daily, weekly, or monthly backup should start, enter how frequently the backup should occur, and enter how many backups to retain.

Example

Consider this example:

- Backup Transfer Start Time: 5 PM
- Transfer Window: 3 hours
- Lag Time: 2 hours

In this example the Snapshot copy transfer will start at 5 PM and will have 3 hours for a transfer window. However, if the backup does not complete in the 3 hours of transfer window, then it has lag time of another 2 hours. By 10 PM, the data should have been transferred.

8. Enter any script information that should be applied either before or after the backup and click **Next**.
9. Review the summary and click **Finish**.

Backing up NAS file services to the cloud

Backing up NAS file data to the cloud includes several short tasks.

About this task

- Back up files to the cloud using SnapCenter and AltaVault
- Monitor data protection progress in the SnapCenter backup Activity pane
- Monitor backup operations progress in the SnapCenter Jobs page

- Verify backup to the cloud completion in AltaVault

Backing up NAS file services to the cloud using SnapCenter and AltaVault

You can back up all shares on a NFS or CIFS protocol-supported volume to the cloud.

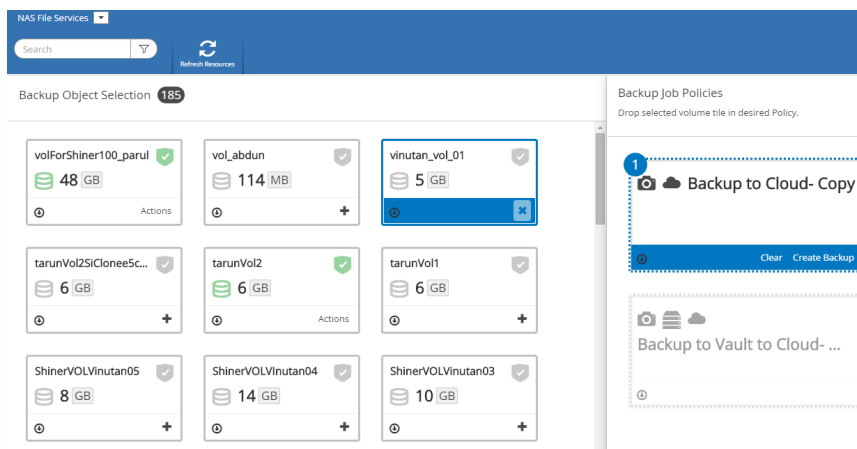
About this task

When you apply a policy to a volume, you can indicate when the initial baseline backup should be performed (either immediately or at a time you specify). After this baseline backup completes, all subsequent backups, which are incremental backups, occur according to the schedule in the policy.

Steps

1. From the SnapCenter left navigation pane, click **Resources**.
2. From the Resources page, search for the volume you want to back up by entering criteria in the Search box and pressing Enter.
3. Optionally, filter the list of results by clicking on the **Filter** icon.

A list of volumes appears.



The green shield on a volume indicates that the volume has been protected already and you cannot set up protection again.

4. To view details about the volume, click the down arrow on the volume.
5. Select a policy among in the policy boxes on the right.

If you do not select a policy, the first policy is selected by default.

Note: Predefined policies do not appear among the policy boxes on the right. You must first copy a predefined policy and modify it before you can use it to govern data protection.

6. To back up the volume to the cloud and index the files within it, drag the volume to the policy box on the right.
7. To view details of the policy and optionally change them, click the down arrow on the policy box. You might want to add tags, for example, "finance," to the backup so that you can later find it by searching for all backups with the tag.
8. To initiate a backup immediately, click the down arrow on the policy box and check **Run initial backup immediately**.

- To finish the backup, in the policy box, click **Create Backup**.

After you finish

Monitor the progress of the data transfer to AltaVault in the Activity pane at the bottom.

Monitoring data protection progress in the SnapCenter Activity pane

After you initiate a backup in SnapCenter, you can immediately review the progress of the data transfer from SnapCenter to AltaVault using the Activity pane that appears on the bottom of the backup page in SnapCenter. You can also track progress for restore operations.


About this task

The Activity pane shows information for the following data protection processes:

- Backup operations to AltaVault
 - Note:** Scheduled backup job activity does not appear on the Activity pane.
- Restore operations from AltaVault if you used the file catalog from the Dashboard or from the Resources option.

Step

- After you initiate a backup or restore operation in SnapCenter, review the SnapCenter Activity pane on the bottom of the page.



The Activity pane shows the following information:

Item	Description
Active, Failed, Completed information in the top bar	Displays a status for all volumes. For example, if two volumes are being backed up at the same time and one failed instantly, the status shows Active = 1 and Failed = 1.
1st column: <number of minutes>	Displays when the activity began.
2nd column: <status>	Displays the activity, such as applying protection on volume, replicating Snapshot, or restoring data.
3rd column: progress bars and icons	<p>The first bar shows the backup progress from ONTAP to AltaVault. The second bar shows the backup progress from ONTAP to ONTAP to AltaVault. During the restore process, the restore activity appears.</p> <ul style="list-style-type: none"> Snapshot icon: Displays the progress of the backup of primary storage data. Vault icon: Displays the progress of the backup from ONTAP to ONTAP to AltaVault. Cloud icon: Displays the progress of the transfer of data from SnapCenter to AltaVault.







Item	Description
4th column: percentage	Shows the percentage complete for the entire data transfer process to AltaVault.

Monitoring backup operations from the Jobs page


You can monitor the progress of different SnapCenter backup operations by using the Jobs page. You might want to check the progress of an operation to determine when it is complete or if there is an issue.


About this task

The following icons appear on the Jobs page and indicate the state of the operation:

-  In progress
-  Completed successfully
-  Failed
-  Completed with warnings or could not start due to warnings
-  Queued
-  Cancelled

Steps

1. In the left navigation pane, click **Monitor**.
2. In the **Monitor** page, click **Jobs**.
3. In the **Jobs** page, perform the following steps:
 - a. Click  to filter the list so that only backup operations are listed.
 - b. Specify the start and end dates.
 - c. From the **Type** drop-down list, select **Backup**.
 - d. From the **Status** drop-down, select the backup status.
 - e. Click **Apply** to view the operations completed successfully.
4. Select the backup job, and then click **Details** to view the job details.

Note: Though the backup job status displays , when you click on job details you might see that some of the child tasks of the backup operation are still in progress.

5. In the **Job Details** page, click **View logs**.

Verifying backup to the cloud completion on AltaVault

A SnapMirror share is created automatically when the SnapMirror relationship with the AltaVault is created in ONTAP. Based on SnapMirror policies created in ONTAP or SnapCenter, Snapshot copies

of ONTAP volumes are backed up to the associated AltaVault share. You should verify that the share was transmitted to AltaVault and then replicated to the cloud.

About this task

Snapshot copies backed up to AltaVault shares are read-only copies and can be restored back to ONTAP using only ONTAP commands or SnapCenter.

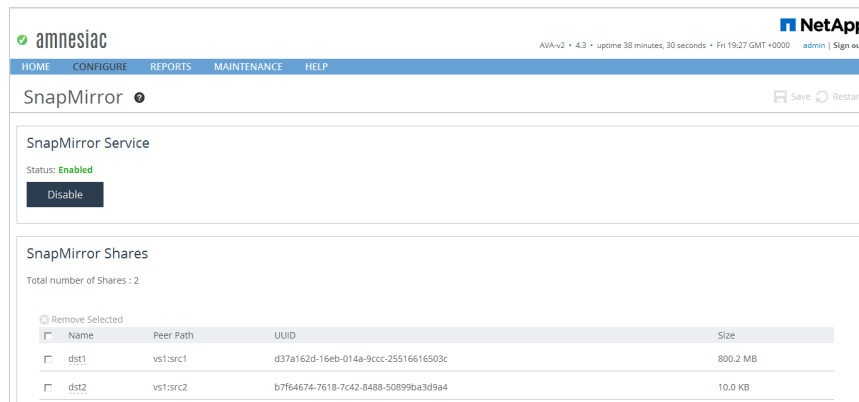
During the lifetime of a share, there is only one baseline Snapshot copy. Any Snapshot copy after the baseline is always incremental. Baseline transfer can take a long time to complete depending on the size of the Snapshot copy.

During incremental backups, only the changed blocks between two Snapshot copies are transferred. Backups can be triggered in ONTAP either through SnapMirror policies in SnapCenter or by explicitly running the ONTAP `snapmirror update` command.

AltaVault provides global deduplication on all Snapshot copy backup streams prior to replication to the cloud.

Steps

1. Log in to AltaVault.
2. Select **Configure > SnapMirror**.
3. Under SnapMirror Shares, review the information associated with share.



Field	Description
<Share> Name	Specifies the name of the share. When the ONTAP administrator creates a SnapMirror relationship with AltaVault or when a NAS file services backup is completed in SnapCenter, a share is automatically created in AltaVault. Each share is associated with one ONTAP FlexVol volume.
Peer Path	Identifies the path to the vServer and source volume in ONTAP (vServer:ONTAP volume) that is being backed up in AltaVault.
UUID	Lists the unique identifier associated each SnapMirror share. The Snapshot copy UUID value is generated by ONTAP. The Share UUID value is generated by AltaVault.

Field	Description
Size	<p>Specifies the size of the SnapMirror share. The size can grow or shrink as Snapshot copies are backed up or deleted from the share.</p> <p>Shares on AltaVault have no size limitation but are bound by the AltaVault appliance capacity. The size of source volume, change rate, and number of Snapshot copies impact the number and size of SnapMirror shares on AltaVault.</p>

- Optionally, select a share and view the Snapshot copy information.

The screenshot shows the NetApp Snapshots interface. At the top, there's a navigation bar with 'HOME', 'CONFIGURE', 'REPORTS', 'MAINTENANCE', and 'HELP'. Below it, the 'Snapshots' section is active. The breadcrumb path is 'SnapMirror Shares > dst1'. A message indicates 'Total number of Snapshots : 2'. A table lists the snapshots:

Name	UUID	Created	Size	Status
snap1	c289835f-e5ef-437d-b3a7-ae22d5bfc6ea	11-11-2016 23:23	10.0 KB	Completed
snap2	8aab267b-04cf-4e43-9db0-e906fd71e48c	11-11-2016 23:42	800.1 MB	Pending

Field	Description
<Snapshot copy> Name	<p>Lists the name of the Snapshot copy. When the ONTAP administrator creates a SnapMirror relationship with AltaVault or when a NAS file services backup is completed in SnapCenter, a share is automatically created in AltaVault. Each share is associated with one ONTAP FlexVol volume.</p> <p>During the lifetime of a share, there is only one baseline Snapshot copy. Any Snapshot copy after the baseline is always incremental. Baseline transfer can take a long time to complete depending on the size of the Snapshot copy. During incremental Snapshot backups, only the changed blocks between two Snapshot copies are transferred.</p>
UUID	Specifies the unique identifier associated each Snapshot copy. The UUID value is generated by ONTAP.
Created	Indicates the date and time when the Snapshot copy was created in ONTAP.
Size	Specifies the size of the Snapshot copy.
Status	<p>Indication of progression of the Snapshot copy backup to the cloud:</p> <ul style="list-style-type: none"> Pending: Snapshot copy replication to the cloud is in progress. Completed: Snapshot copy replication to the cloud was successful. <p>Snapshots appear only after they have been ingested into AltaVault.</p>

- Optionally, to remove a share or Snapshot copy, select the item from the list and click **Remove Selected**.

Snapshot copies can be deleted on AltaVault through ONTAP or SnapCenter policies, or by manual deletion on AltaVault. Upon Snapshot copy deletion, the data belonging to a Snapshot copy is not deleted immediately from AltaVault. AltaVault has an asynchronous reclamation process, meaning the data is reclaimed only when a certain number of Snapshot copies have been deleted.

When a share is deleted, Snapshot copies in that share are also deleted. Snapshot copies cannot be restored once deleted even though the Snapshot copy data may not have been deleted.

Verifying backup to the cloud completion using AltaVault CLI

Based on SnapMirror policies created in ONTAP or SnapCenter, Snapshot copies of ONTAP volumes are backed up to the associated AltaVault share. You might want to verify that the share was transmitted to AltaVault and then replicated to the cloud. You can perform this using the AltaVault command-line interface (CLI).

Steps

- Use SSH to log in to the AltaVault CLI and enter configuration mode:

```
<hostname> > enable
<hostname> # configuration terminal
<hostname> (config) #
```

- After the backup has occurred, look at the SnapMirror shares on AltaVault:

```
<hostname> (config)# show snapmirror shares
Total Snapmirror shares - 2

Share: dst1
  Peer path: vs1:src1
    UUID: a7f871fb-b757-cb42-86fc-1b3cf8746346
    Size: 20.0 KB
Share: dst2
  Peer path: vs1:src2
    UUID: 6aa603ba-0084-1248-a71c-6de5dbbdac92
    Size: 20.0 KB
```

- After the backup has occurred, look at the Snapshot copies on those shares on AltaVault:

```
<hostname> (config)# show snapmirror snapshots share-name dst1
Total snapshots - 2

Name:
  UUID:                               Create Time:      Size:
Replication Status
snap1
  82c5671a-746b-4358-86a7-50744a0bea33  11-11-2016 13:11  10.0 KB      Complete
snap2
  2d825301-33bc-4aee-a8fa-c75230a57ddb  11-11-2016 13:11  10.0 KB      Pending
```

Restoring NAS file data from the cloud

You can restore the entire volume or a single file from the cloud to ONTAP. Restoring NAS file data from the cloud includes several short tasks.

About this task

- Restore files from the cloud using SnapCenter and AltaVault
- Monitor data protection progress in the SnapCenter backup Activity pane

- Monitor restore operations progress in the SnapCenter Jobs page

As a best practice, we recommend that you initiate no more than 10 concurrent restore operations. Otherwise, transfer time outs and slower restore throughput may occur. SnapCenter supports up to 100 concurrent transfers to a single AltaVault node, including both backup and restore operations. This means that if 100 concurrent backups are running against a single node, you cannot initiate any restore operations from that node.

Searching in the file catalog for a specific file to restore using keywords and filters

Using the powerful SnapCenter file catalog, you can search for a specific file to restore among thousands of files across volumes. You can narrow down the list of results quickly and easily.

About this task

You can search for NAS file services file that you want to restore in the following ways:

- From the SnapCenter Dashboard: If you know the specific file, keywords in that file name, type of file, or date of the file, but do not know the volume that stores the file, you should use the file catalog.
For example, you can find a specific file from years ago that you want to restore. You can filter by volume or share name, tag names assigned to backups, file types, file size, owner, and data range.
- From the SnapCenter Resources page: If you know the volume name, you should search for volumes using the Resource page. Then, you can filter by volume, storage virtual machine (SVM), share, or path.
For details, see information about restoring volumes.

Cataloging of files is performed only in the primary storage system. The file catalog is updated with information on where the backup exists (primary, secondary, or AltaVault) after the data transfer is complete. Indexing of files for the catalog is done only at the file level, not on the contents inside files.

If there is at least one share in a volume, the SnapCenter file catalog includes only the shares. Otherwise, the catalog includes the entire volume.

Steps

1. From the SnapCenter Dashboard, click **Restore Backups**.
2. Search for the file you want to restore using the following search tips:
 - Search for single terms.
 - To search by multiple words, surround the phrase with double quotes.
 - To combine multiple phrases, use Boolean operators in all capital letters: AND, OR, NOT.

```
SnapMirror AND "archival mode"
```

 - To search for any single character, enter "?" for that character, even in the middle of a term, for example:


```
te?t
```
 - To search for multiple characters, enter "*" at the end or even in the middle of the term. For example, "test*" yields the results of "test," "testing," and "tester."

Note: You cannot use a * or ? as the first character of a search.

- To search for similar "fuzzy" terms, enter the tilde "~" at the end of the term. For example, "backup~" could yield terms like "rack" and "backed."

```
backup~
```

- To search for words that are within a specific number of words away from each other, use the tilde "~" at the end of a phrase. For example, to search for "licensing" and "SnapMirror" within 10 words of each other in the document, enter the following phrase that is surrounded by double quotes:

```
"licensing SnapMirror"~10
```

- To include a special character such as the following in your query, use the \ character before the special character:

```
+ - && | | ! ( ) { } [ ] ^ " ~ ? : \
```

A list of files appears, which could include possibly thousands of files that match your search criteria.

- To find the specific file you want to restore, filter the results by volume or share name, tag names assigned to backups, file types, file size, owner, or data range.

Filters

Share Name

Volume

From

mm/dd/yyyy --:-- --

To

mm/dd/yyyy --:-- --

Owner

Tag

finance

File Types

e.g. *.txt, *.doc

File Size

e.g. > 5 MB

Clear Apply

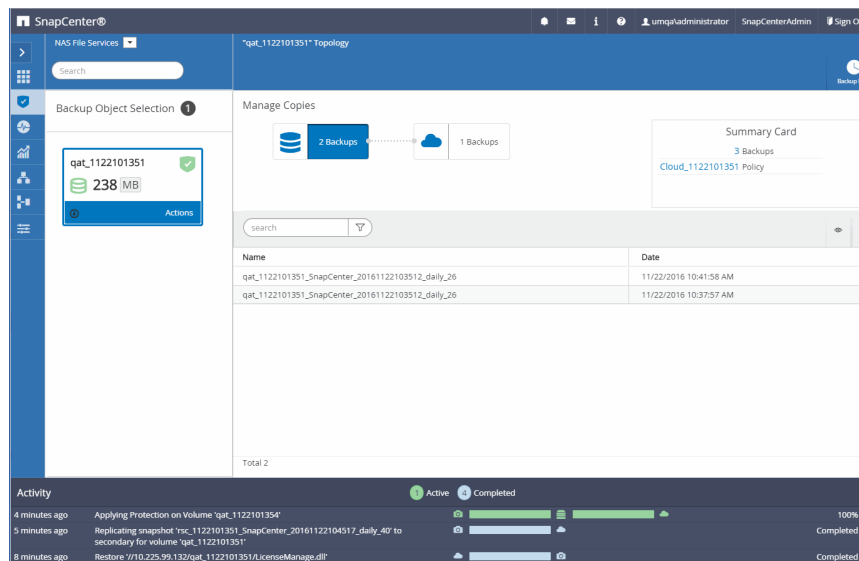
- Click **Apply**.
- From the list of filtered results, click on the file you want to restore.

Restoring NAS file data from the cloud using SnapCenter and AltaVault

When you want to restore a NAS file that has been damaged, accidentally changed or deleted, you can use SnapCenter to restore the file from a cloud backup to its original location.

Steps

1. In the left navigation pane, click **Resources**.
2. In the **Resources** page, locate and select the volume in the **Backup Object Selection** view from which you want to restore a file, and then click **Actions**.
3. In the **Manage Copies** view, select the backup you want to use to restore a file, and then click the **Browse** (eye) icon.
4. Expand the folders as needed to locate the file you want to restore, and then select the file.
5. In the **Restore Files** view, select to restore it to its original location:
6. Click **Restore**.



7. In the confirmation message that all data written to the file will be overwritten with the backup data, click **Yes** to continue.
8. Monitor the progress of the operation by viewing the **Activity** pane at the bottom of the page.
The Activity pane shows information for the following data protection processes:
 - Backup operations to AltaVault
 - Restore operations from AltaVault if you used the file catalog from the Dashboard or if you used the Resources option.
9. Monitor the progress of the operation using the **Jobs** page.

Related tasks

[Monitoring restore operations from the Jobs page](#) on page 57

[Monitoring data protection progress in the SnapCenter Activity pane](#) on page 49







Monitoring restore operations from the Jobs page

You can monitor the progress of different SnapCenter restore operations by using the Jobs page. You might want to check the progress of an operation to determine when it is complete or if there is an issue.


About this task

Post-restore states describe the conditions of the resource after a restore operation and any further restore actions that you can take.

The following icons appear on the Jobs page, and indicate the state of the operation:

-  In progress
-  Completed successfully
-  Failed
-  Completed with warnings or could not start due to warnings
-  Queued
-  Cancelled

Steps

1. In the left navigation pane, click **Monitor**.
2. In the **Monitor** page, click **Jobs**.
3. In the **Jobs** page, perform the following steps:
 - a. Click  to filter the list so that only restore operations are listed.
 - b. Optional: Specify the start and end dates.
 - c. From the **Type** drop-down list, select **Restore**.
 - d. From the **Status** drop-down list, select the restore status.
 - e. Click **Apply** to view the operations that are completed successfully.
4. Select the restore job, and then click **Details** to view the job details.
5. In the **Job Details** page, click **View logs**.

Managing SnapMirror shares using the AltaVault CLI

Rather than use the AltaVault graphical user interface, you can use the AltaVault command-line interface (CLI) for managing SnapMirror shares.

About this task

If you use the AltaVault graphical user interface, you do not need to perform this procedure.

SnapMirror shares are created automatically after a SnapMirror relationship to the AltaVault is created either in ONTAP or in SnapCenter.

Steps

1. Use SSH to log in to the AltaVault CLI:

```
<hostname> > enable
```

2. Display the list of shares created on the AltaVault:

```
<hostname> (config) # show snapmirror shares
```

To delete a share from the AltaVault, use the `snapmirror share delete sharename <share-name>` command.

3. Display the list of all Snapshot copies for a share:

```
<hostname> (config) # show snapmirror share <share-name> snapshots
```

AltaVault displays the Snapshot copy name, UUID, size, and replication status to the cloud. Replication status can be Completed or Pending.

To delete a Snapshot copy from a share, use the `snapmirror snapshots delete share-name <share-name> snapshot-name name <snapshot name>` command

Backing up NAS file services to the cloud and restoring them using ONTAP commands

Using the ONTAP command-line interface (CLI), you can back up NAS file services to the cloud and restore them from the cloud back to an ONTAP volume.

About this task

Instead of using ONTAP commands, you can initiate backups using the SnapCenter graphical user interface (GUI). See information about backing up NAS file services using SnapCenter.

One ONTAP volume, which can include multiple directories and files, can be backed up to one share in AltaVault. First, you back up a full baseline to AltaVault and then you can perform incremental backups forever.

You can later restore a full backup including all Snapshot copies from AltaVault to a new destination ONTAP volume or you can restore a single file in a Snapshot copy from AltaVault to a new or existing ONTAP volume.

Backing up NAS file services to the cloud using ONTAP commands

Backing up NAS file services to the cloud using ONTAP command-line interface (CLI) involves performing a baseline backup first followed by incremental backups. ONTAP FlexVols are backed up to AltaVault as NAS file shares on volumes.

Related tasks

[Backing up NAS file services to the cloud and restoring them with SnapCenter management](#) on page 45

Performing a baseline backup from ONTAP to the cloud using the ONTAP CLI

When you perform a baseline backup, you review the volume's contents, create a SnapMirror relationship, create a Snapshot copy, initialize the SnapMirror relationship, transfer the data to AltaVault, and finally show the transfer of the contents in AltaVault.

About this task

A baseline backup creates a SnapMirror relationship with an AltaVault endpoint.

Steps

1. Log in to the ONTAP CLI.
2. Look at the contents of the primary volume:

Example

```
cluster1::> volume show
Vserver Volume      Aggregate State  Type    Size    Available Used%
-----
vs2      ontap_source aggr1    online RW     1GB     966.4MB 5%
vs2      vs_root1    aggr1    online RW     20MB     18.01MB 9%
3 entries were displayed.
```

3. Look at the files on the volume:

Example

```
cluster1::> set diag
cluster1::> run local ls -lr /vol/ontap_source
```

4. Create a SnapMirror relationship, which creates the share on AltaVault:

For a backup of an ONTAP FlexVol to an ONTAP FlexVol, both the source and destinations use the format <Vserver>:<volume>. However for the AltaVault endpoint, you use the format <IP_address>:/share/<share_name>

The protection type must be "XDP" for Extended Data Protection, which is often used for SnapVault protection.

```
cluster1::> snapmirror create
-source-path <Vserver><volume>:
-destination-path <IP_address>:/share/<share_name> -type XDP
```

Example

```
cluster1::> snapmirror create
-source-path vs2:ontap_source
-destination-path 192.0.2.150:/share/ava_dest_1 -type XDP
```

5. Verify that the SnapMirror relationship is established and is in the Unitialized state:

To view the detailed status of the relationship, use the -instance option.

Example

```
cluster1::> destination_cluster::>snapmirror show
Source      Destination Mirror      Relationship Total      Progress
Path  Type Path      State      Status      Progress Healthy Updated
-----
vs34:src    XDP  192.0.2.22:/share/dst Uninitialized Idle -   true   -
1 entries were displayed.

destination_cluster::> snapmirror show -instance
Source Path: vs34:src
Destination Path: 172.19.201.22:/share/dst
Relationship Type: XDP
Relationship Group Type: none
SnapMirror Schedule: -
SnapMirror Policy Type: vault
SnapMirror Policy: XDPDefault
```

6. Create a Snapshot copy:

Example

```
cluster1::> snapshot create -vserver vs2 -volume ontap_source -snapshot base_snapshot
```

7. Initialize the SnapMirror relationship and transfer the data to AltaVault:

```
cluster1::> snapmirror initialize
-source-path <Vserver>:<volume_name>
-destination-path <IP_address>:/share/<share_name>
-source-snapshot <Snapshot_name>
```

Example

```
cluster1::> snapmirror initialize -s
-source-path vs2:ontap_source
-destination-path 192.0.2.150:/share/ava_dest1
-source-snapshot base_snapshot
```

If you use `Snapmirror initialize` without the `-s` (explicit snapshot), the Snapshot copy is deleted in the next update.

You can initialize the relationship only once during the lifetime of the relationship.

8. Show the transfer of the contents:

Example

```
cluster1::> snapmirror show
```

Source Path	Type	Destination Path	Mirror State	Relationship Status	Total Progress	Progress Healthy	Last Updated
vs34:src	XDP	192.0.2.150:/share/dst	Snapmirrored	Idle	-	true	

The `Idle` state indicates the transfer is complete.

9. Show the transfer of the contents:

Example

```
cluster1::> snapshot show ontap_source
```

Vserver	Volume	Snapshot	Size	---Blocks---	Total%	Used%
vs2	ontap_source	base_snapshot	108KB	0%	2%	
		incremental_snapshot	92KB	0%	1%	
		hourly.2016-10-22_0905	252KB	0%	4%	
		hourly.2016-10-22_1005	92KB	0%	1%	
		hourly.2016-10-22_1105	100KB	0%	2%	

10. Additionally, using the AltaVault CLI, look for the baseline backup of shares on the destination volume in AltaVault.

Example

```
onk-sm# show snapmirror shares

onk-sm# show snapmirror snapshots share-name ava_dest_1
Name:
UUID          Create Time:      Replication Status
base_snapshot
8477ad3d-95bb-4398-9da8-358b5d79cf45
10-22-2016 15:46   Completed
Total Snapshots - 1
```

The Replication Status changes from `Pending` to `Completed`, indicating that the replication to AltaVault has completed.

After you finish

Perform an incremental backup from ONTAP to AltaVault.

Performing an incremental backup from ONTAP to the cloud using the ONTAP CLI

After completing a baseline backup, you should perform incremental backups that capture additional Snapshot copy data since the baseline backup occurred.

Before you begin

You must have already completed a baseline backup from ONTAP to AltaVault.

About this task

When you perform an incremental backup from ONTAP to AltaVault and to the cloud, you review the volume's contents, create a SnapMirror relationship, create a Snapshot copy, update the SnapMirror relationship, transfer the data to AltaVault, and finally show the transfer of the contents in AltaVault.

Steps

1. Log in to the ONTAP CLI.
2. Create another Snapshot copy:

Example

```
cluster1::> snapshot create -vserver vs2
               -volume ontap_source
               -snapshot base_snapshot
```

Alternatively, you can schedule a backup by using the command: `snapmirror policy create` or schedule the update by using the command: `snapmirror modify -destination-path <dest-path> -schedule <schedule>`

3. Create an incremental backup:

Example

```
cluster1::> snapmirror update
               -source-path vs2:ontap_source
               -destination-path 192.0.2.150:/share/ava_dest_1
               -source-snapshot incremental_snapshot
```

Do not use the `-enable-storage-efficiency` option with the `snapmirror update` command for Data Fabric Solution for Cloud Backup operations.

4. Show the transfer of the contents:

Example

```
cluster1::> snapmirror show
```

Source Path	Type	Destination Path	Mirror State	Relationship Status	Total Progress	Progress Last Updated	Healthy
vs2:ontap_source	XDP	192.0.2.150:/share/ava_dest_1	mirrored	Transferring	0B	10/22	true

11:28:00

The Relationship Status changes from Transferring to Idle, indicating that the transfer is complete.

5. Show the transfer of the contents:

Example

```
cluster1::> snapmirror show
```

Source Path	Destination Type Path	Mirror State	Relationship Status	Total Progress	Progress Last	Healthy Updated
vs2:ontap_source	XDP	192.0.2.150:/share/dst	Snapmirrored	Idle	-	true

6. Additionally, using the AltaVault CLI, look for baseline and incremental backup of shares on the destination volume in AltaVault.

Example

```
onk-sm# show snapmirror snapshots share-name ava_dest_1
```

Name:	UUID	Create Time:	Replication Status
base_snapshot	8477ad3d-95bb-4398-9da8-358b5d79cf45	10-22-2016 15:46	Completed
incremental_snapshot	41a9268a-67c9-40bb-8c2c-76c5c02fc597	10-22-2016 15:58	Completed
Total Snapshots	-	2	

The Replication Status changes from Pending to Completed, indicating that the replication to AltaVault has completed.

After you finish

Using AltaVault, configure the cloud provider settings. After these are set up, the backup of data that was transferred to AltaVault is automatically replicated to the cloud. See AltaVault documentation.

Monitoring backup progress using the ONTAP CLI

After you initiate a backup, you can monitor the progress of that backup to the cloud using the ONTAP command-line interface (CLIL).

Steps

1. Log in to the ONTAP CLI.
2. Show the progress of the Snapshot copy transfer:

```
cluster1::>snapmirror show
```

Example

```
cluster1::*> snapmirror show
```

Source Path	Destination Type Path	Mirror State	Relationship Status	Total Progress	Progress Last	Healthy Updated
vs34:src	XDP	192.0.2.22:/share/dst	Uninitialized	Transferring	12.86MB true	10/24

```
cluster1::*> snapmirror show -instance 192.0.2.22:/share/dst
```

```

Source Path: vs34:src
Source Cluster: -
Source Vserver: vs34
Source Volume: src
Destination Path: 192.0.2.22:/share/dst
Destination Cluster: -
Destination Vserver: vs34
Destination Volume: -
Relationship Type: XDP
Relationship Group Type: none
Managing Vserver: vs34
SnapMirror Schedule: -
SnapMirror Policy Type: vault

```

```

SnapMirror Policy: XDPDefault
Tries Limit: -
Throttle (KB/sec): unlimited
Current Transfer Throttle (KB/sec): 0
Mirror State: Uninitialized
Relationship Status: Transferring
File Restore File Count: -
File Restore File List: -
Transfer Snapshot:
snapmirror.b8f64f8f-94cb-11e6-8193-0050568576ed_1101672341869.2016-10-24_123309
Snapshot Progress: 12.86MB
Total Progress: 12.86MB
Network Compression Ratio: 1:1
Snapshot Checkpoint: 0B
...

```

3. Show the progress of the Snapshot copy transfer again until you see that the Mirror State shows "Snapmirrored" and the Relationship status shows "Idle" as in the following example:

```
cluster1::>snapmirror show
```

Example

```
cluster1::> snapmirror show
```

Source Path	Type	Destination Path	Mirror State	Relationship Status	Total Progress	Healthy	Progress Last Updated
vs34:src	XDP	192.0.2.22:/share/dst	Snapmirrored	Idle	-	true	-

Restoring NAS file services from the cloud using ONTAP commands

You can restore a full backup from the cloud to ONTAP or you can restore a single file from a Snapshot copy in AltaVault back to ONTAP.

About this task

You can restore data in the following ways:

- Restore a full backup from the AltaVault cloud to a new destination volume in ONTAP
- Restore a single file from a Snapshot copy in AltaVault to a new or existing volume in ONTAP

As a best practice, we recommend that you initiate no more than 10 concurrent restore operations. Otherwise, transfer time outs and slower restore throughput may occur. SnapCenter supports up to 100 concurrent transfers to a single AltaVault node, including both backup and restore operations. This means that if 100 concurrent backups are running against a single node, you cannot initiate any restore operations from that node.

Restoring a full backup from the cloud using the ONTAP CLI

Restoring a full baseline backup from the cloud to ONTAP includes creating a new destination volume in ONTAP into which you will restore the shares, restoring AltaVault data back to the new destination volume in ONTAP, reviewing the transfer of data, and ensuring that the ONTAP volume reverts to a read/write volume.

Before you begin

You must have already backed up data from ONTAP to AltaVault.

About this task

You can only perform a baseline full volume restore of any Snapshot copy backed up to AltaVault. If you have say, for example, 365 incremental restore operations, you would restore 365 volumes separately.

Alternatively, you can restore a single file to ONTAP rather than the entire baseline backup.

Steps

1. Log into the AltaVault CLI.
2. Using the AltaVault CLI, identify the share:

```
<hostname> (config) # show snapmirror shares
```

3. Using the AltaVault CLI, identify the Snapshot copies for a share:

```
<hostname> (config) # show snapmirror snapshots share-name <share_name>
```

4. Log in to the ONTAP CLI.
5. Using the ONTAP CLI, create a destination volume where you will restore the Snapshot data using a relationship type of DP:

```
cluster1::> volume create -volume <volume> -type DP
```

Example

The following command creates a 2 GB volume of type DP:

```
cluster1::> volume create -volume ontap_dest
-aggregate aggl -type DP -size 2GB
```

For a single file restore, the destination volume must be read/write. For a full volume restore, the transfer of a share from AltaVault back to ONTAP converts the volume automatically to a read/write volume.

6. Look at the list of volumes to see the volume you created:

Example

```
cluster1::> volume show
```

7. Restore AltaVault share data back to the ONTAP volume:

```
cluster1::> Snapmirror restore
-source-path <IP_address>:/share/<share_name>
-destination-path <Vserver_name>:<volume_name>
-source-snapshot <Snapshot_name>
```

Do not use the `snapmirror restore -use-network-compression` option for Data Fabric Solution for Cloud Backup.

Example

```
cluster1::> Snapmirror restore
-source-path 192.0.2.150:/share/ava_dest_1
-destination-path vs2:ontap_dest
-source-snapshot base_snapshot
```

8. Look at the data transferred from AltaVault back to ONTAP:

Example

```
cluster1::> Snapmirror show
```

Source Path	Destination Type Path	Mirror State	Relationship Status	Total Progress	Progress Healthy	Last Updated
192.0.2.150:/share/ava_dest_1	RST	vs2:ontap_dest	Unitialized	Transferring	0B	
vs2:ontap_source	XDP	192.0.2.150:/share/ava_dest_1	Snapmirrored	Idle	true	

2 entries were displayed.

The Mirror State of RST indicates that restore relationships exist while the transfer is occurring. After the transfer completes, the RST relationship is removed.

- Again, look at the data transferred from AltaVault back to ONTAP:

Example

```
cluster1::> Snapmirror show
```

Source Path	Destination Type Path	Mirror State	Relationship Status	Total Progress	Progress Healthy	Last Updated
192.0.2.150:/share/ava_dest_1	RST	vs2:ontap_dest	Snapmirrored	Idle	true	
vs2:ontap_source	XDP	192.0.2.150:/share/ava_dest_1	Snapmirrored	Idle	true	

2 entries were displayed.

The Relationship Status changes from Unitialized Transferring to Snapmirrored.

- Again, look at the data transferred from AltaVault back to ONTAP:

Example

```
cluster1::> Snapmirror show
```

Source Path	Destination Type Path	Mirror State	Relationship Status	Total Progress	Progress Healthy	Last Updated
192.0.2.150:/share/ava_dest_1	RST	vs2:ontap_dest	Broken-off	Idle	true	
vs2:ontap_source	XDP	192.0.2.150:/share/ava_dest_1	Snapmirrored	Idle	true	

2 entries were displayed.

The Relationship Status changes from Snapmirrored to Broken-off indicating that the relationship has been removed.

- Look at the volume on ONTAP and verify that the destination relationship changes from DP to R/W:

Example

```
cluster1::> volume show
```

Vserver	Volume	Aggregate	State	Type	Size	Available	Used%
onk-vs1ml	vol0	aggr0	online	RW	2.85GB	705.4MB	76%
vs2	ontap_dest	aggr1	online	RW	1GB	1021MB	0%
vs2	ontap_source	aggr1	online	RW	1GB	966.4MB	5%
vs2	vs_root1	aggr1	online	RW	20MB	10MB	9%

4 entries were displayed.

- Verify that the Snapshot copy was restored on ONTAP:

Example

```
cluster1::> snapshot show ontap_dest
```

Vserver	Volume	Snapshot	Size	---Blocks---	Total%	Used%
vs2	ontap_dest	base_snapshot	112KB	0%	5%	

Related tasks

[Monitoring backup progress using the ONTAP CLI](#) on page 63

Restoring a single file from the cloud to ONTAP using the ONTAP CLI

Instead of restoring an entire volume, you can restore individual files from the Snapshot copy back to an existing or new ONTAP volume. Restoring a single file from the cloud to ONTAP includes creating a new destination volume if needed in ONTAP into which you will restore the file, restoring AltaVault data back to the destination volume in ONTAP, reviewing the transfer of data, and ensuring that the ONTAP volume reverts to a read/write volume.

Before you begin

You must have already backed up data from ONTAP to AltaVault.

Steps

1. Log into the AltaVault CLI.
2. Using the AltaVault CLI, identify the share:
3. Using the AltaVault CLI, identify the Snapshot copies for a share:

```
<hostname> (config) # show snapmirror shares
```

```
<hostname> (config) # show snapmirror snapshots share-name <share_name>
```

4. Log in to the ONTAP CLI.
5. Using the ONTAP CLI, optionally create a destination volume where you will restore the Snapshot data using a relationship type of DP:

```
cluster1::> volume create -volume <volume> -type RW -size <volume_size>
```

Example

The following command creates a 1 GB volume named "sfr_dest" of type Read/Write:

```
cluster1::> volume create -volume sfr_dest
-aggregate aggl -type RW -size 1GB -state online
```

You can restore a single file back to a new or existing volume in ONTAP. If the file that you are restoring exists in the existing volume, the existing file will be overwritten.

For a single file restore, the destination volume must be read/write. For a full volume restore, the transfer of a share from AltaVault back to ONTAP converts the volume automatically to a read/write volume.

6. Look at the list of volumes to see the volume you created:

Example

```
cluster1::> volume show
```

7. Restore the single file from AltaVault to the ONTAP volume:

```
cluster1::> Snapmirror restore
-source-path <IP_address>:/share/<share_name>
-destination-path <Vserver_name>:<volume_name>
-source-snapshot <snapshot_name>
-file-list /<filename>
```

Do not use the `snapmirror restore -use-network-compression` option for Data Fabric Solution for Cloud Backup.

Example

This example restores a single file "file3" from the Snapshot copy in AltaVault named "incremental_snapshot" to an ONTAP destination volume named "sfr_dest" :

```
cluster1::> Snapmirror restore
-source-path 192.0.2.150:/share/ava_dest_1
-destination-path vs2:sfr_dest
-source-snapshot incremental_snapshot
-file-list /file3
```

8. Look at the data transferred from AltaVault back to ONTAP:

Example

```
cluster1::> snapmirror show
```

Source Path	Destination Type Path	Mirror State	Relationship Status	Total Progress	Progress Last Updated	Healthy
192.0.2.150:/share/ava_dest_1		RST	vs2:sfr_dest			
			Idle			true
vs2:ontap_source	XDP	192.0.2.150:/share/ava_dest_1				
		Snapmirrored	Idle			true

2 entries were displayed.

The Mirror State of RST indicates restore.

9. List the contents of the volume to verify that file was restored:

Example

```
cluster1::> set diag
cluster1::> run local ls -lr /vol/sfr_dest
10644 96 1 2097152 Fri Oct 21 12:02:59 EST 2016 file3
```

Related tasks

[Monitoring backup progress using the ONTAP CLI](#) on page 63

Viewing data protection reports for NAS file services

Using the Dashboard page, you can view reports about backup and restore operations for NAS file services that have a particular status. This is useful if you want to identify the total number of successful or failed operations in your SnapCenter environment.

About this task

Report type	Description
Backup Report	The Backup Report provides overall data about backup trends for NAS file services, the backup success rate, and some information about each backup performed during the specified time. If a backup is deleted, the report does not display any status information for the deleted backup. The Backup Detail Report provides detailed information about a specified backup job and lists the resources successfully backed up and any that have failed.
Restore Report	The Restore Report provides overall information about restore operations. The Restore Detail Report provides details about a specified restore operation, including host name, backup name, job start and duration, and the status of individual job tasks. If a task fails, the Restore Detail Report displays information about the failure.
Data Protection Report	This report provides protection details for NAS file services. You can see how many volumes are or are not protected.

Steps

1. In the left navigation pane, click **Dashboard**.
2. Locate the data protection pie chart.
3. Click the pie slice representing the status for which you want a report.
You are directed to the Reports page. The report displays only operations with the status you selected.
4. Review the report, download it to your local system, print it, export it, or have it emailed.
To email the report, select the **Settings > Global Settings** option and configure the SMTP server.

Where to go next

You can find more information about different features and release-specific information for Data Fabric Solution for Cloud Backup in the documentation available on the NetApp Support Site at mysupport.netapp.com.

- ONTAP Documentation Center at <http://docs.netapp.com/ontap-9/index.jsp>
- *NetApp AltaVault Cloud Integrated Storage Administration Guide*
NetApp Documentation: AltaVault
- *SnapCenter Software 2.0 Release Notes*
Provides important information about this release of SnapCenter Server and the SnapCenter plug-in packages, including fixed issues, known issues, cautions, limitations, and any documentation updates or corrections.
- *SnapCenter Software 2.0 Installation and Setup Guide*
Describes the steps required to prepare for installation and to install SnapCenter and the SnapCenter plug-in packages. Setup processes are described for both current SnapManager users who are importing data to SnapCenter and users who are implementing a new SnapCenter environment.
- *SnapCenter Software 2.0 Administration Guide*
Provides information about how to administer SnapCenter, provision Windows hosts with storage, configure and maintain role-based access control (RBAC), and use the centralized reporting options.
- *SnapCenter Software 2.0 Windows Cmdlet Reference Guide*
Provides reference information about the Windows PowerShell cmdlets available in SnapCenter, including a description of each cmdlet, its syntax, and examples for its use. This content is also available through the SnapCenter PowerShell cmdlet help.
- *SnapCenter Software 2.0 Linux Command Reference Guide*
Provides reference information about the Linux commands available for Linux plug-ins, including a description of each command, its syntax, and examples for its use. This content is also available through the SnapCenter command-line interface help.

Copyright information

Copyright © 1994–2017 NetApp, Inc. All rights reserved. Printed in the U.S.

No part of this document covered by copyright may be reproduced in any form or by any means—graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system—without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark information

Active IQ, AltaVault, Arch Design, ASUP, AutoSupport, Campaign Express, Clustered Data ONTAP, Customer Fitness, Data ONTAP, DataMotion, Element, Fitness, Flash Accel, Flash Cache, Flash Pool, FlexArray, FlexCache, FlexClone, FlexPod, FlexScale, FlexShare, FlexVol, FPolicy, Fueled by SolidFire, GetSuccessful, Helix Design, LockVault, Manage ONTAP, MetroCluster, MultiStore, NetApp, NetApp Insight, OnCommand, ONTAP, ONTAPI, RAID DP, RAID-TEC, SANscreen, SANshare, SANtricity, SecureShare, Simplicity, Simulate ONTAP, Snap Creator, SnapCenter, SnapCopy, SnapDrive, SnapIntegrator, SnapLock, SnapManager, SnapMirror, SnapMover, SnapProtect, SnapRestore, Snapshot, SnapValidator, SnapVault, SolidFire, SolidFire Helix, StorageGRID, SyncMirror, Tech OnTap, Unbound Cloud, and WAFL and other names are trademarks or registered trademarks of NetApp, Inc., in the United States, and/or other countries. All other brands or products are trademarks or registered trademarks of their respective holders and should be treated as such. A current list of NetApp trademarks is available on the web.

<http://www.netapp.com/us/legal/netapptmlist.aspx>

How to send comments about documentation and receive update notifications

You can help us to improve the quality of our documentation by sending us your feedback. You can receive automatic notification when production-level (GA/FCS) documentation is initially released or important changes are made to existing production-level documents.

If you have suggestions for improving this document, send us your comments by email.

doccomments@netapp.com

To help us direct your comments to the correct division, include in the subject line the product name, version, and operating system.

If you want to be notified automatically when production-level documentation is released or important changes are made to existing production-level documents, follow Twitter account @NetAppDoc.

You can also contact us in the following ways:

- NetApp, Inc., 495 East Java Drive, Sunnyvale, CA 94089 U.S.
- Telephone: +1 (408) 822-6000
- Fax: +1 (408) 822-4501
- Support telephone: +1 (888) 463-8277

Index

A

- access
 - assigning to users [36](#)
 - enabling SnapCenter in AltaVault [30](#)
 - logging in to AltaVault [27](#)
 - logging in to SnapCenter [34](#)
- administrator
 - logging in to SnapCenter as [34](#)
- administrator role
 - adding in AltaVault [30](#)
- AltaVault
 - adding a SnapCenter user [30](#)
 - configuring [27](#)
 - configuring whitelist of approved IP addresses [28](#)
 - configuring whitelist of approved IP addresses in AltaVault CLI [31](#)
 - enabling access to SnapCenter [30](#)
 - enabling SnapMirror long-term retention mode [29](#)
 - installation [26](#)
 - logging in [27](#)
 - managing shares with the AltaVault CLI [57](#)
 - managing SnapMirror shares [50](#)
 - managing SnapMirror shares using CLI [53](#)
 - starting Storage Optimization Service [31](#)
 - supported versions [11](#)
- architecture
 - components [6](#)
 - prerequisites [9](#)
- ASM instance
 - setting up your Run As account credentials [39](#)
- authentication
 - setting up your Run As account [39](#)
- AutoSupport notifications
 - enabling for storage system connections [37](#)

B

- backing up
 - monitoring NAS file services backup using ONTAP CLI [63](#)
 - monitoring progress [49](#)
 - monitoring progress in SnapCenter [50](#)
 - NAS file services using SnapCenter GUI [47](#), [48](#)
 - performing a baseline backup of NAS file services using ONTAP CLI [59](#)
 - performing an incremental backup of NAS file services using ONTAP CLI [62](#)
 - using ONTAP CLI, overview [59](#)
 - using or creating policies [45](#)
- backing up NAS file services
 - with ONTAP CLI, overview [59](#)
- best practices
 - concurrent restore operations [45](#), [53](#), [64](#)
 - setting up your Run As account [39](#)
 - SnapCenter file catalog [10](#)
- browsers
 - supported versions [12](#)

C

- catalog
 - Linux server availability [23](#)
 - protecting [41](#), [43](#)
 - requirements [10](#)
 - searching for files to restore [54](#)
- CIFS
 - configuring server [41](#)
 - required licenses for SnapCenter [12](#)
- cloud backup
 - components [6](#)
 - overview [5](#)
- clusters
 - creating cluster peering relationship [19](#)
 - credential requirements [15](#)
- comments
 - how to send feedback about documentation [73](#)
- configuring
 - AltaVault overview [27](#)
 - CIFS server [41](#)
 - intercluster LIFs [19](#)
 - SnapCenter overview [34](#)
 - workflow diagram with SnapCenter implementation [16](#), [17](#)
- connections
 - required before installing [15](#)
 - setting up for storage systems [37](#)
- credentials
 - requirements [15](#)
 - setting up your Run As account [39](#)

D

- Data Fabric Solution for Cloud Backup
 - connection and port requirements for installation [15](#)
 - documentation resources [70](#)
 - features [6](#)
 - installation order [18](#)
 - installing AltaVault [26](#)
 - installing ONTAP [18](#)
 - overview [5](#)
 - preparing for deployment [9](#)
 - what you can do [8](#)
- data transfer
 - monitoring cloud backup completion in AltaVault [50](#)
 - to AltaVault [49](#)
- databases
 - preparing the PowerShell environment [42](#)
- documentation
 - additional resources [70](#)
 - how to receive automatic notification of changes to [73](#)
 - how to send feedback about [73](#)

E

- EMS messaging

enabling for storage systems connections [37](#)

F

- feedback
 - how to send comments about documentation [73](#)
- file catalog
 - Linux server availability [23](#)
 - protecting [43](#)
 - requirements [10](#)
 - searching for files to restore [54](#)
- FlexClone licenses
 - required for SnapCenter [12](#)

H

- hosts
 - licenses for SnapCenter plug-ins [12](#)
 - operating systems supported [10](#)
 - space and sizing requirements for installing Data Fabric Solution for Cloud Backup [10](#)
 - supported systems [9](#)
 - system requirements for installing Data Fabric Solution for Cloud Backup [10](#)
- HTTP
 - port requirements [15](#)

I

- information
 - how to send feedback about improving documentation [73](#)
- installing
 - AltaVault [26](#)
 - connection and port requirements [15](#)
 - ONTAP [18](#)
 - operating systems supported [10](#)
 - order of components [18](#)
 - prerequisites [9](#)
 - SnapCenter [24](#)
 - space and sizing requirements for packages [10](#)
 - space requirements for packages [10](#)
 - workflow diagram with SnapCenter implementation [16, 17](#)
- intercluster LIFs
 - configuring [19](#)

J

- jobs
 - monitoring backup progress from the Jobs page [50](#)
 - monitoring restore progress from the Jobs page [57](#)

L

- licenses
 - adding for ONTAP [22](#)
 - for Data Fabric Solution for Cloud Backup [12](#)
- Linux
 - setting up your Run As account credentials on host [39](#)

- Linux host
 - ensuring accessibility [23](#)
- logging in
 - AltaVault [27](#)
 - as a SnapCenter administrator [34](#)
 - as a SnapCenter user [34](#)
 - SnapCenter [34](#)
 - with more than one role [34](#)

M

- monitoring
 - cloud backup completion in AltaVault [50](#)
 - data protection progress in SnapCenter Activity pane [49](#)
 - progress of SnapCenter backup operations [50](#)
 - progress of SnapCenter restore operations [57](#)

N

- NAS file services
 - backing up to the cloud [48](#)
 - backing up to the cloud overview [45, 47](#)
 - cloud backup overview [5](#)
 - configuring AltaVault [27](#)
 - configuring SnapCenter [34](#)
 - protecting file catalog [43](#)
 - protecting indexing catalog [41](#)
 - reporting from the Dashboard [69](#)
 - restoring from the cloud [56](#)
 - restoring from the cloud overview [45, 53](#)
 - searching in catalog for a file to restore [54](#)
- NFS
 - required licenses for SnapCenter [12](#)
- non-root user
 - installing plug-in for Oracle databases [23](#)

O

- ONTAP
 - adding a license [22](#)
 - installation [18](#)
 - supported versions [9](#)
- ONTAP CLI
 - backing up and restoring NAS file services overview [59](#)
 - backing up NAS file services to the cloud, overview [59](#)
 - monitoring progress of cloud backup [63](#)
 - performing a baseline backup of NAS file services [59](#)
 - performing an incremental backup of NAS file services [62](#)
 - restoring a single file from the cloud to ONTAP [67](#)
 - restoring from the cloud to ONTAP, overview [64](#)
 - restoring NAS file services from the cloud to ONTAP [64](#)
- operating systems
 - supported [10](#)
- operations
 - monitoring backup progress from the Jobs page [50](#)
 - monitoring restore progress from the Jobs page [57](#)

Oracle Database
 setting up your Run As account credentials [39](#)

P

permissions
 setting for applications [39](#)
 plug-ins
 space and sizing requirements for Data Fabric
 Solution for Cloud Backup installation [10](#)
 policies
 using or creating [45](#)
 port requirements
 before installing [15](#)
 PowerShell
 preparing the environment [42](#)
 protecting
 NAS file services file catalog [43](#)
 SnapCenter metadata repository [43](#)

R

reports
 data protection [69](#)
 repository
 protecting [41](#), [43](#)
 requirements
 connection and port, before installing Data Fabric
 Solution for Cloud Backup [15](#)
 credentials [15](#)
 restore operations
 monitoring progress from the Jobs page [57](#)
 restoring
 full baseline restore using ONTAP CLI [64](#)
 searching in catalog for a file [54](#)
 single file from the cloud [56](#)
 single file from the cloud, overview [53](#)
 single file restore using ONTAP CLI [67](#)
 using ONTAP CLI [64](#)
 restoring NAS file services
 with ONTAP CLI, overview [59](#)
 role-based access control
 adding account in AltaVault [30](#)
 adding users to a role in SnapCenter [35](#)
 assigning access to users [36](#)
 Run As account
 setting up [39](#)

S

services
 Storage Optimization [31](#)
 shares
 managing in AltaVault [50](#)
 managing in AltaVault CLI [53](#)
 SnapCenter
 configuring [34](#)
 enabling access to in AltaVault [30](#)
 finding the URL [34](#)
 installing [24](#)
 installing overview [23](#)
 SnapMirror

configuring with the AltaVault CLI [31](#)
 enabling long-term retention mode in AltaVault [29](#)
 licenses required for SnapCenter [12](#)
 managing shares with the AltaVault CLI [57](#)

SnapMirror shares
 managing in AltaVault [50](#)
 managing using AltaVault CLI [53](#)
 managing with the AltaVault CLI [57](#)
 SnapRestore
 licenses required for SnapCenter [12](#)
 Snapshot copies
 enabling AltaVault access to SnapCenter [30](#)
 impact on retention by SnapMirror long-term
 retention mode [29](#)
 SnapVault
 licenses required for SnapCenter [12](#)
 SQL Server
 setting up your Run As account credentials [39](#)
 storage system connections
 enabling AutoSupport notifications for [37](#)
 enabling EMS messaging for [37](#)
 storage systems
 licenses for SnapCenter plug-ins [12](#)
 settings up connections to [37](#)
 sudo privileges
 configuring for non-root user [23](#), [39](#)
 suggestions
 how to send feedback about documentation [73](#)
 systems
 requirements for Data Fabric Solution for Cloud
 Backup installation [10](#)
 space and sizing requirements for SnapCenter
 installation [10](#)
 supported for Data Fabric Solution for Cloud
 Backup [10](#)

T

Twitter
 how to receive automatic notification of
 documentation changes [73](#)

U

users
 adding a SnapCenter user in AltaVault GUI [30](#)
 adding to a role in SnapCenter [35](#)
 assigning access to [36](#)
 using roles at login [34](#)

W

web address
 for SnapCenter [34](#)
 whitelist of approved IP addresses
 configuring in AltaVault CLI [31](#)
 configuring in AltaVault GUI [28](#)
 Windows Server
 operating systems supported [10](#)
 workflows
 preparing the PowerShell environment [42](#)