



ONTAP® 9

Volume Move Express Guide

September 2016 | 215-11197_B0
doccomments@netapp.com

Visit the new ONTAP 9 Documentation Center: docs.netapp.com/ontap-9/index.jsp

Contents

Deciding whether to use this guide	4
Volume move workflow	5
Planning the method and timing of a volume move	6
Moving a volume using System Manager	7
Moving a volume using the CLI (advanced)	8
Selecting the destination aggregate	9
Adding LUN reporting nodes before moving a volume	11
Starting the volume move	11
Monitoring the volume move	12
Performing the volume move cutover	13
Updating LUN reporting nodes after moving a volume	13
Updating NDMP backup after moving a volume	14
Where to find additional information	15
Copyright information	16
Trademark information	17
How to send comments about documentation and receive update	
notifications	18
Index	19

Deciding whether to use the Volume Move Express Guide

You can use this guide to nondisruptively move a volume from one node to another node within the same Storage Virtual Machine (SVM) in an ONTAP 9.0 cluster.

Requirements for using this guide

Before you use this guide, ensure that the following conditions are met:

- The cluster is running Data ONTAP 9.0.
- You have cluster administrator privileges.
- You want to use best practices, not explore every available option.
- You do not want to read a lot of conceptual background.
- You can use either OnCommand System Manager or the Data ONTAP command-line interface.
- You know which volume you want to move.
For help deciding which volumes to move, you can use OnCommand Performance Manager.
- The volume that will be moved is a data volume, not a root volume.
- Any new or repurposed hardware is fully installed and already has aggregates.
- If the cluster has LUNs, all nodes have two paths per LUN.
- Flow control is not enabled on cluster network ports.

If this guide is not suitable for your situation, you should see the following documentation instead:

- [Logical storage management](#)
- [NetApp Knowledgebase](#)
The Knowledgebase contains instructions on moving the root volume.
- [NetApp Documentation: OnCommand Workflow Automation \(current releases\)](#)
OnCommand Workflow Automation enables you to run prepackaged workflows that automate management tasks such as the workflows described in Express Guides.

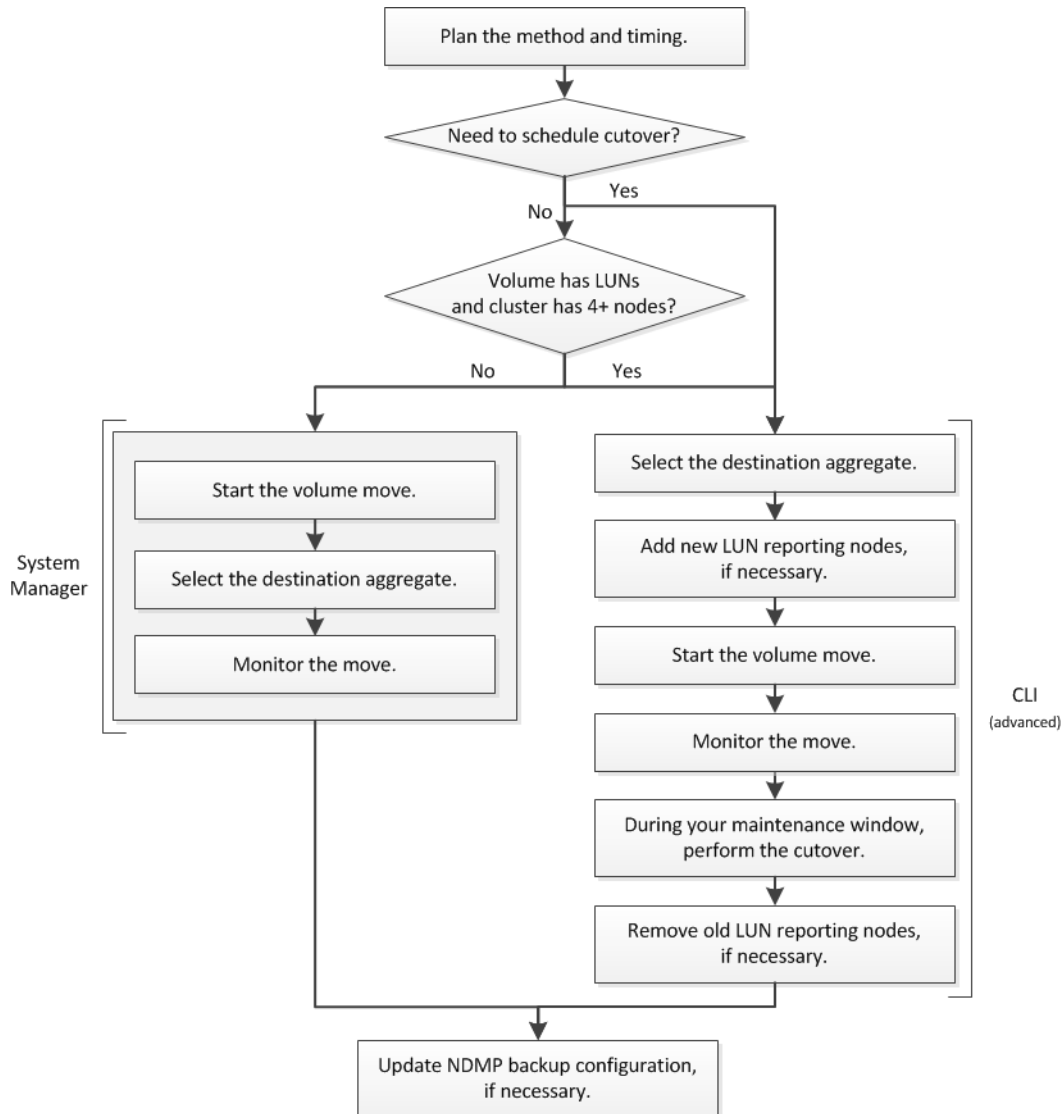
Alternatives to volume move

Before moving volumes, you should evaluate whether the following approaches are better suited to your situation:

- If you want to nondisruptively upgrade a controller in place, you can consider using aggregate relocation (ARL), which does not require physical data movement.
[High-availability configuration](#)
- If you want to move only a LUN—without its containing volume—you can use the LUN move mechanism.
[SAN administration](#)

Volume move workflow

Before moving a volume, you should select a method and plan the timing. In most cases, you can use the simple System Manager procedure. If your situation requires it, you should use the more advanced CLI method. After the move, you might have to update the NDMP backup configuration.



Planning the method and timing of a volume move

In most situations, you can use System Manager to move a volume. If you need to control the final cutover or update LUN reporting nodes, you must follow an advanced procedure in the command-line interface (CLI). You can also optionally plan the timing of a volume move.

About this task

The source volume of a SnapMirror or SnapVault relationship can be moved while the volume is being mirrored. SnapMirror services encounter a brief pause during the cutover phase of the volume move job.

The destination volume can also be moved. In the iterative phase, SnapMirror or SnapVault updates and volume move operations run concurrently. When evaluating whether a cutover is possible in the cutover phase, priority between the cutover and SnapMirror or SnapVault updates is determined on a first-come, first-served basis. The first operation to run blocks other operations until the first operation finishes.

Steps

1. Decide whether you have to use the CLI because you require a manual cutover.

Cutover is the moment at which the move operation finishes and Data ONTAP starts serving data from the volume on the new aggregate. The cutover can occur automatically or you can trigger it manually.

- If your company's standard practice requires you to control when changes occur in the storage system, use the CLI to manually perform the final cutover of the move operation during a maintenance window.

A cutover does not require an outage, but you can use a maintenance window to control *when* it occurs.

- If it is acceptable for the cutover to occur automatically, use System Manager. Although you can use the CLI to move a volume with automatic cutover, it is easier to use System Manager.

Note: The volume move operation is nondisruptive, regardless of whether you choose automatic or manual cutover.

2. Decide whether you have to use the CLI because you must update LUN reporting nodes:

- If the volume does not contain LUNs or if the cluster contains only two nodes, use System Manager.
- If the volume contains LUNs and the cluster contains four or more nodes, use the CLI so that you can update LUN reporting nodes if the volume moves to a different HA pair.

3. Optional: Plan a time using the following considerations:

- A volume move operation might take more time than expected because moves are designed to occur nondisruptively in the background in a manner that preserves client access and overall system performance.

For example, Data ONTAP throttles the resources available to the volume move operation.

- If you want the move to occur as quickly as possible, select a time with less cluster activity, especially the following activities:
 - I/O operations on the volume

- Jobs using background resources, for example, when controller CPU usage is less than 50 percent
- Jobs using the cluster interconnect
- A move cannot be started while the volume is affected by the following operations: volume offline, restrict, or destroy; SnapMirror resync, break, or restore; and Snapshot restore. You must wait for any of these specific operations to finish before you can start the move.
- While the volume move operation occurs, a MetroCluster switchback cannot occur, although a switchover can occur.
- MetroCluster switchbacks are blocked when volume move operations are in progress for volumes belonging to the switched over site. Switchbacks are not blocked when volume move operations are in progress for volumes local to the surviving site.
- Forced MetroCluster switchovers can occur when volume move operations are in progress.

Moving a volume using System Manager

You can easily move a volume in System Manager by selecting a volume, starting the move, and then optionally monitoring the resulting job. When using System Manager, a volume move operation finishes automatically.

Steps

1. Select the volume:
 - a. Expand the **Storage Virtual Machines** hierarchy in the left navigation pane.
 - b. Select **Storage > Volumes**.
 - c. In the **Volumes** window, select the volume you want to move.
 - d. Click **Move**.

The Move Volume dialog box is displayed.

Move Volume

Source Volume

Name: vol1
Committed Size: 204.8 MB
Aggregate: aggr1
Storage Type: SSD

Destination Aggregate

Name	Available Space	Total Space	RAID Type	Storage Type
aggr3	763.8 MB	784.35 MB	raid_dp, normal	SSD

Source Aggregate Available Space

Before Move: 763.8 MB After Move: 968.6 MB

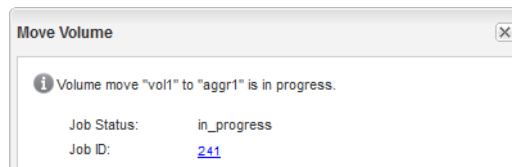
Destination Aggregate Available Space

Before Move: 763.8 MB After Move: 559 MB

Move Cancel

2. Select the destination aggregate and start the move:

- a. Review the list of possible aggregates, which includes only the aggregates that have the required capacity.
You should review the available space, total space, RAID type, and storage type. For example, if the goal is to alter performance characteristics of the volume, you can focus on aggregates with the desired storage type.
- b. Select a destination aggregate from the list.
- c. Review the available space on the source and destination aggregates before the move and after the move.
- d. Click **Move**, and then click **Move** again to confirm the action.
- e. When a new **Move Volume** dialog box is displayed, leave it open so that you can use it for the next step.



3. Optional: Monitor the volume move job:

- a. In the new **Move Volume** dialog box, click the link to the **Job ID** of the volume move job.

If the dialog box is already closed, you can expand the **Cluster** hierarchy, and then select **Diagnostics > Jobs**.

The Jobs window is displayed, with the newest jobs at the bottom of the list.

- b. Locate the Volume Move job, and review the **Status** column.

The job can be in any one of several phases, such as transferring the initial baseline of data or starting a cutover attempt.

241	03/05/2015 07:3...	Volume Move	node1-1	running	Move "vol1" in V...	Cutover Started:(1 of 3 attempts) Transferring final da...
-----	--------------------	-------------	---------	---------	---------------------	--

- c. Continue to monitor the job either by remaining on the **Jobs** window and clicking **Refresh** or by periodically returning to the **Jobs** window.

241	03/05/2015 07:3...	Volume Move	node1-1	success	Move "vol1" in V...	Complete: Successful [0]
-----	--------------------	-------------	---------	---------	---------------------	--------------------------

The job status is `Complete: Successful` when the volume move finishes.

Moving a volume using the CLI (advanced)

When using the CLI to move a volume within a maintenance window or with updates to LUN reporting nodes, you select a destination aggregate, optionally update LUN reporting nodes, and start

the move. After monitoring the operation until it is ready, you then trigger the cutover and optionally update LUN reporting nodes.

Steps

1. [Selecting the destination aggregate](#) on page 9
2. [Adding LUN reporting nodes before moving a volume](#) on page 11
3. [Starting the volume move](#) on page 11
4. [Monitoring the volume move](#) on page 12
5. [Performing the volume move cutover](#) on page 13
6. [Updating LUN reporting nodes after moving a volume](#) on page 13

Selecting the destination aggregate

Before you start moving a volume through the command-line interface, you must identify the aggregate to which you will move the volume. The `volume move target-aggr show` command identifies all possible destination aggregates for you to evaluate for your specific situation.

Steps

1. Identify the volume's SVM and its current aggregate (known as the *source* aggregate) by using the `volume show` command with the `-volume` and `-instance` parameters.

Example

```
cluster1::> volume show -volume voll -instance
Vserver Name: SVM1
Volume Name: voll
Aggregate Name: aggr1
...
```

2. Display all aggregates that are potential destinations for this volume by using the `volume move target-aggr show` command.

The output includes only data aggregates with enough available capacity for the volume.

Example

```
cluster1::> volume move target-aggr show -vserver SVM1 -volume voll
Aggregate Name  Available Size  Storage Type
-----
aggr3           763.8GB        SSD
aggr4           345.5GB        SSD
```

3. Select a possible destination aggregate from the list.
4. Ensure that the destination aggregate is assigned to the SVM:
 - a. Determine whether the SVM has aggregates assigned to it by using the `vserver show` command with the `-vserver` and `-fields aggr-list` parameters.

If the list of aggregates is empty, aggregate assignment is not a factor in choosing a destination aggregate.

Example

The following output displays an SVM with an empty list of aggregates:

```
cluster1::> vserver show -vserver SVM1 -fields aggr-list
vserver aggr-list
-----
SVM1      -
```

- b. If the `aggr-list` contains aggregates and the destination aggregate is not among them, select a different aggregate or use the `vserver add-aggregates` command to assign more aggregates to the SVM.
5. Ensure that the destination aggregate is on the desired node:
 - a. Identify the node of the source aggregate by using the `storage aggregate show` command with the `-aggregate` and `-fields node` parameters.

Example

```
cluster1::> storage aggregate show -aggregate aggr1 -fields node
aggregate node
-----
aggr1      node-1-1
```

- b. Identify the node of the destination aggregate by using the `storage aggregate show` command with the `-aggregate` and `-fields node` parameters.

Example

```
cluster1::> storage aggregate show -aggregate aggr3 -fields node
aggregate node
-----
aggr3      node-1-2
```

- c. If the aggregate is not on the desired node—for example, if the destination aggregate is on the same node as the source aggregate—select a different aggregate.
6. If the cluster contains more than one disk type, ensure that the destination aggregate is of the desired disk type.
 - a. Identify the storage type of the source aggregate by using the `storage aggregate show` command with the `-aggregate` and `-fields aggregate-type` parameters.

Example

```
cluster1::> storage aggregate show -aggregate aggr1 -fields aggregate-type
aggregate storage-type
-----
aggr1      SSD
```

- b. Identify the disk type of the destination aggregate by using the `storage aggregate show` command with the `-fields aggregate-type` parameter.

Example

```
cluster1::> storage aggregate show -aggregate aggr3 -fields aggregate-type
aggregate storage-type
-----
aggr3      SSD
```

- c. If the aggregate does not have the desired storage type, select a different aggregate.

For example, if the goal is to alter performance characteristics of the volume, you must select the aggregate with the desired disk type.

Adding LUN reporting nodes before moving a volume

If the volume that you plan to move contains LUNs and the destination aggregate is on another HA pair, you should add the new HA pair to the Selective LUN Map (SLM) reporting-nodes list. This ensures that optimized LUN paths are maintained after the volume is moved.

About this task

This procedure is necessary only if you plan to move the volume to a different HA pair. If the volume is on a different node of the same HA pair—for example, you have a two-node cluster or a MetroCluster configuration—you can skip this procedure.

Steps

1. Add the destination node and its partner node to the volume's reporting-nodes list by using the `lun mapping add-reporting-nodes` command with the `-destination-aggregate` parameter set to the aggregate that you selected as the destination.

Example

```
cluster1::> lun mapping add-reporting-nodes -vserver SVM1 -volume voll -lun * -igroup ig1
               -destination-aggregate aggr3
```

2. Verify that the LUN map contains four nodes by using the `lun mapping show -fields reporting-nodes` command.

Example

```
cluster1::> lun mapping show -vserver SVM1 -volume voll -fields reporting-nodes
vserver  path      igroup  reporting-nodes
-----
SVM1     /vol/voll  ig1     cluster1-1,cluster1-2,cluster1-3,cluster1-4
```

If the cluster contains only four nodes and all four are included as reporting-nodes, the `reporting-nodes` field displays a dash (-).

3. Rescan from the host to discover the newly added paths.
4. Add the new paths to your MPIO configuration.

Starting the volume move

You can start the process of moving a volume to another aggregate in a way that defers the final cutover phase until a maintenance window of your choice.

Before you begin

If the volume is a data protection mirror, the mirror relationship must be initialized.

Steps

1. Verify that the volume move is permitted by the system by using the `volume move start` command with the `-perform-validation-only` parameter.

The command verifies that no current operations prevent the start of the volume move operation. For example, it is not possible to move a volume that is undergoing a SnapMirror resync or Snapshot restore operation. If mutually exclusive operations are running, you must wait for them to finish before continuing.

2. Set the privilege level to advanced by using the `set -privilege advanced` command.

Advanced privilege is required so that you can use the `-cutover-action` parameter that defers cutover until it is manually triggered. Without the `-cutover-action` parameter, the cutover occurs automatically.

3. Start the volume move operation by using the `volume move start` command with the `-cutover-action wait` parameter.

Example

```
cluster1::> volume move start -vserver SVM1 -volume voll -destination-aggregate aggr3 -
cutover-action wait

Warning: When the transfer is complete, trigger the cutover to avoid loss of
space savings from efficiency. Do you want to proceed?
{y|n}: y
[Job 153] Job is queued: Move "voll" in Vserver "SVM1" to aggregate "aggr3". Use the
"volume move show -vserver SVM1 -volume voll" command to view the status of this
operation.
```

4. Return the privilege level to admin by using the `set -privilege admin` command.

Monitoring the volume move

After starting a volume move operation, you should use the `volume move show` command to monitor its progress and identify when the move is ready for manual cutover.

Steps

1. View information about the volume move operation by using the `volume move show` command.

After an **initializing** phase, the move operation enters a **replicating** phase.

Example

```
cluster1::> volume move show
Vserver  Volume  State  Move Phase Percent-Complete Time-To-Complete
-----
SVM1     voll    healthy replicating 62% Tue Mar 10 15:33:16 20
```

2. Monitor the progress of the Percent Complete field by repeatedly using the `volume move show` command.

Example

```
cluster1::> volume move show
Vserver  Volume  State  Move Phase Percent-Complete Time-To-Complete
-----
SVM1     voll    healthy replicating 93% Tue Mar 10 15:33:16 20
```

Tip: You can use the date and time in the Time-To-Complete field as an indicator of when to monitor the operation.

3. When the Move Phase changes to **cutover_hard_deferred**, schedule the cutover at a convenient time, such as a maintenance window.

Example

```
cluster1::> volume move show
Vserver  Volume  State  Move Phase Percent-Complete Time-To-Complete
-----
SVM1     voll    alert  cutover_hard_deferred - -
```

The move operation is ready for you to manually trigger the cutover. In the meantime, any changes made to the volume are periodically replicated to the destination aggregate.

Performing the volume move cutover

When the volume move operation is in a deferred cutover state and you are ready to finish the volume move operation, you can manually trigger the cutover to the new aggregate.

Steps

1. Finish the volume move operation by using the `volume move trigger-cutover` command.

Example

```
cluster1::> volume move trigger-cutover -vserver SVM1 -volume voll
```

After the current transfer is complete, cutover begins.

2. Verify that the cutover has started by using the `volume move show` command.

Example

The following output shows an operation that is in the **cutover** phase and then the **finishing** phase.

```
cluster1::> volume move show -vserver SVM1 -volume voll
Vserver  Volume  State  Move Phase Percent-Complete Time-To-Complete
-----
SVM1     voll    healthy cutover    -              -
```

```
cluster1::> volume move show -vserver SVM1 -volume voll
Vserver  Volume  State  Move Phase Percent-Complete Time-To-Complete
-----
SVM1     voll    healthy finishing  -              -
```

3. Verify that the entire operation is finished by repeatedly using the `volume move show` command with the `-vserver`, `-volume`, and `-instance` parameters until the state is **done**.

Example

The following output shows an operation that is 100% complete, in a **completed** move phase, and in a **done** state.

```
cluster1::> volume move show -vserver SVM1 -volume voll -instance
Vserver Name: SVM1
Volume Name: voll
...
Detailed Status: Successful
Percentage Complete: 100%
...
Move Phase: completed
Move State: done
```

If the operation fails, no action is required; the volume remains on the source aggregate. You can review the `Detailed Status` field for information about the cause.

Updating LUN reporting nodes after moving a volume

If the volume that you moved contained LUNs and the volume is now on a different HA pair, you should remove all remote nodes from the Selective LUN Map (SLM) reporting-nodes list. The LUN

map then contains only the owner node and its HA partner, which ensures that only optimized LUN paths are used.

About this task

This procedure is necessary only if you moved the volume to a different HA pair. If the volume is on a different node of the same HA pair, you can skip this procedure.

Steps

1. Remove all remote nodes from the reporting-nodes list by using the `lun mapping remove-reporting-nodes` command with the `-remote-nodes` parameter.

Example

```
cluster1::> lun mapping remove-reporting-nodes -vserver SVM1 -volume voll -igroup ig1 -remote-nodes true
```

2. Verify that the LUN map contains only the owner node and its partner by using the `lun mapping show` command with the `-fields reporting-nodes` parameter.

Example

```
cluster1::> lun mapping show -vserver SVM1 -volume voll -fields reporting-nodes
vserver  path      igroup  reporting-nodes
-----
SVM1     /vol/voll  ig1     cluster1-3,cluster1-4
```

3. Remove stale device entries for the host operating system.
4. Rescan from the host to refresh the host's available paths.

See your host documentation for specific steps to rescan your hosts.

Updating NDMP backup after moving a volume

If the volume that you moved was previously backed up to tape using NDMP in a specific configuration, after moving the volume, you can perform one of the following actions to ensure the volume continues to be backed up successfully: create a baseline *or* migrate the backup LIF to the node containing the moved volume.

About this task

- This procedure is necessary only if the backup application does not support the cluster-aware backup (CAB) extension and the backup process uses node-scoped NDMP.
If the backup application supports CAB and it is configured to use SVM-scoped NDMP mode, you can skip this procedure.
- You must perform only one of these actions, not both.

Choices

- From the backup application, create a new baseline.
- Identify the LIF that is configured for the backup process, and then migrate the LIF to the node where the volume now resides.

Where to find additional information

If you want more information about administering ONTAP, you can see the ONTAP 9 product library.

Related information

[NetApp Documentation: Data ONTAP 8 \(current releases\)](#)

Copyright information

Copyright © 1994–2016 NetApp, Inc. All rights reserved. Printed in the U.S.

No part of this document covered by copyright may be reproduced in any form or by any means—graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system—without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark information

NetApp, the NetApp logo, Go Further, Faster, AltaVault, ASUP, AutoSupport, Campaign Express, Cloud ONTAP, Clustered Data ONTAP, Customer Fitness, Data ONTAP, DataMotion, Fitness, Flash Accel, Flash Cache, Flash Pool, FlashRay, FlexArray, FlexCache, FlexClone, FlexPod, FlexScale, FlexShare, FlexVol, FPolicy, GetSuccessful, LockVault, Manage ONTAP, Mars, MetroCluster, MultiStore, NetApp Insight, OnCommand, ONTAP, ONTAPI, RAID DP, RAID-TEC, SANtricity, SecureShare, Simplicity, Simulate ONTAP, Snap Creator, SnapCenter, SnapCopy, SnapDrive, SnapIntegrator, SnapLock, SnapManager, SnapMirror, SnapMover, SnapProtect, SnapRestore, Snapshot, SnapValidator, SnapVault, StorageGRID, Tech OnTap, Unbound Cloud, and WAFL and other names are trademarks or registered trademarks of NetApp, Inc., in the United States, and/or other countries. All other brands or products are trademarks or registered trademarks of their respective holders and should be treated as such. A current list of NetApp trademarks is available on the web.

<http://www.netapp.com/us/legal/netapptmlist.aspx>

How to send comments about documentation and receive update notifications

You can help us to improve the quality of our documentation by sending us your feedback. You can receive automatic notification when production-level (GA/FCS) documentation is initially released or important changes are made to existing production-level documents.

If you have suggestions for improving this document, send us your comments by email.

doccomments@netapp.com

To help us direct your comments to the correct division, include in the subject line the product name, version, and operating system.

If you want to be notified automatically when production-level documentation is released or important changes are made to existing production-level documents, follow Twitter account @NetAppDoc.

You can also contact us in the following ways:

- NetApp, Inc., 495 East Java Drive, Sunnyvale, CA 94089 U.S.
- Telephone: +1 (408) 822-6000
- Fax: +1 (408) 822-4501
- Support telephone: +1 (888) 463-8277

Index

A

- about this guide
 - deciding whether to use the Volume Move Express Guide [4](#)
- aggregates
 - destination for moving volume [9](#)
- audience
 - for the volume move express guide [4](#)

B

- backup
 - NDMP, updating after volume move [14](#)

C

- CLI
 - moving volumes with [8](#)
- command-line interface
 - See* CLI
- comments
 - how to send feedback about documentation [18](#)
- cutover
 - deciding between manual and automatic when moving volumes [6](#)
 - deferring when starting a volume move [11](#)
 - performing manually [13](#)

D

- destination
 - aggregates, selecting [9](#)
- documentation
 - how to receive automatic notification of changes to [18](#)
 - how to send feedback about [18](#)
- documentation, express guide
 - where to find additional information about administering ONTAP [15](#)

E

- express guides
 - requirements for using the volume move guide [4](#)
 - volume move workflow [5](#)
 - where to find additional information about administering ONTAP [15](#)

F

- feedback
 - how to send comments about documentation [18](#)
- FlexVol volume move
 - See* volume move
- flowcharts
 - volume move workflow [5](#)

I

- information
 - how to send feedback about improving documentation [18](#)

L

- LUN map
 - See* Selective LUN Mapping
- LUNs
 - adding reporting node before volume move [11](#)

M

- monitoring
 - volume move using System Manager [7](#)
 - volume move using the CLI [12](#)
- moving
 - volumes using System Manager [7](#)
 - volumes using the CLI [8](#)
 - volumes, workflow for [5](#)
 - with deferred cutover [11](#)

N

- NDMP backup
 - updating after volume move [14](#)
- nodes
 - adding LUN reporting node before volume move [11](#)
- nondisruptive volume move
 - volumes
 - moving, workflow for [5](#)

O

- operations, volume move
 - planning the method and timing of [6](#)

P

- planning
 - destination of volume move [9](#)
 - timing and method of volume move [6](#)

S

- Selective LUN Mapping
 - adding nodes before volume move [11](#)
 - removing nodes after volume move [13](#)
- SLM
 - See* Selective LUN Mapping
- suggestions
 - how to send feedback about documentation [18](#)
- System Manager
 - moving volumes with [7](#)

T

target

See destination

timing

deciding for volume move operation [6](#)

Twitter

how to receive automatic notification of
documentation changes [18](#)

V

vol move

See volume move

volume move

adding LUN reporting nodes using the CLI [11](#)

cutover using the CLI [13](#)

monitoring using System Manager [7](#)

monitoring using the CLI [12](#)

removing LUN reporting nodes using the CLI [13](#)

selecting destination aggregate using the CLI [9](#)

starting through the CLI [11](#)

updating NDMP backup after [14](#)

using the CLI [8](#)

workflow for [5](#)

volume move operations

planning the method and timing of [6](#)

volumes

moving using System Manager [7](#)

moving using the CLI [8](#)

requirements for using this guide to move [4](#)

W

workflows

volume move [5](#)